


# Hive to Iceberg

## To Migrate, or Not to Migrate

Starburst Webinar  
May 8, 2024





# Connection Before Content

**Lester Martin** - <https://about.me/lestermartin>

- Educational Engineer @ Starburst
  - Build the content
  - Teach the class
  - Repeat
- 30 years of technology experience
  - Started my journey on a TRS-80 Model III
  - Played most every role, but consider myself a programmer at my core
  - Half of career in transactional systems and the second half in analytical processing
  - A DECADE of “big data” experience to include
    - Trino/Starburst, Hadoop, Hive, Spark
    - NiFi, Kafka, Storm, Flink
    - HBase, MongoDB

# Webinar Agenda

## Slides, but DEMOs, too!

- Evolution of a data lakehouse  
*(the 3 min version)*
- Picking your components
- Building a data lakehouse
- When NOT to migrate from Hive
- Migration strategies
- Additional considerations

Scan for a Trino and Iceberg cheat sheet



# Evolution of the data lakehouse

How did we get here?

# Data Architecture Evolution

## Data Warehouse




**Charmander**

## Data Lake



**Charmeleon**

## Data Lakehouse



**Charizard**

# The Data Warehouse



Popularized in the 90's to provide a 360 degree view

## The Good

- Integrates siloed RDBMS's into one "centralized" location
- Simple & reliable analytical querying
- Data audit, governance and lineage
- Great for small amounts of data

## The Bad

- Inability to store unstructured data
- Lack scalability and flexibility
- Tightly coupled storage and compute
- Expensive, proprietary hardware and software (*creating vendor lock-in*)

# The Data Lake



Born out of the internet age and big data boom


## The Good

- In 2006, Apache Hadoop emerges so unstructured data can be processed at a scale previously unimaginable
- Shift toward parallel processing
- Capitalize on low cost object storage
- Allows for greater flexibility (schema on read)

## The Bad

- Inability to support transactions, updates, or modifications
- Difficult to get top-tier performance
- Lack of data quality and inconsistent data formats
- Insufficient data lineage and limited data discoverability

# The Data Lakehouse



## Applying data warehouse principles to the data lake

- Utilize the ***separation of storage and compute*** to apply the reliability, performance, and data quality of the data warehouse to the openness and scalability of the data lake
- ***Increased performance and scalability*** through the use of indexing and caching via your query engine (**Trino**) and modern table formats (ex: **Iceberg**)
- Provide traditional ***data modifications*** (ex: `UPDATE` & `MERGE` commands) with ***ACID transaction*** guarantees over files stored in the data lake
- Tackle ***unstructured, semi-structured, and structured*** analytical data all in a data lakehouse - creating a place for AI/ML & BI use cases alike





# Picking your components

Trino is the best query engine ever

# The data accessibility problem

## Data practitioners faced the same challenges at Facebook in 2010

- Facebook created Hive to query terabytes of data in Hadoop using SQL
- Data scientists attempted to query massive object stores, but performance was too slow
- Data consumers were limited by the number of queries they could run — often ***fewer than 10*** in one day




# Enter Trino (Presto)

A new open source query engine designed for speed

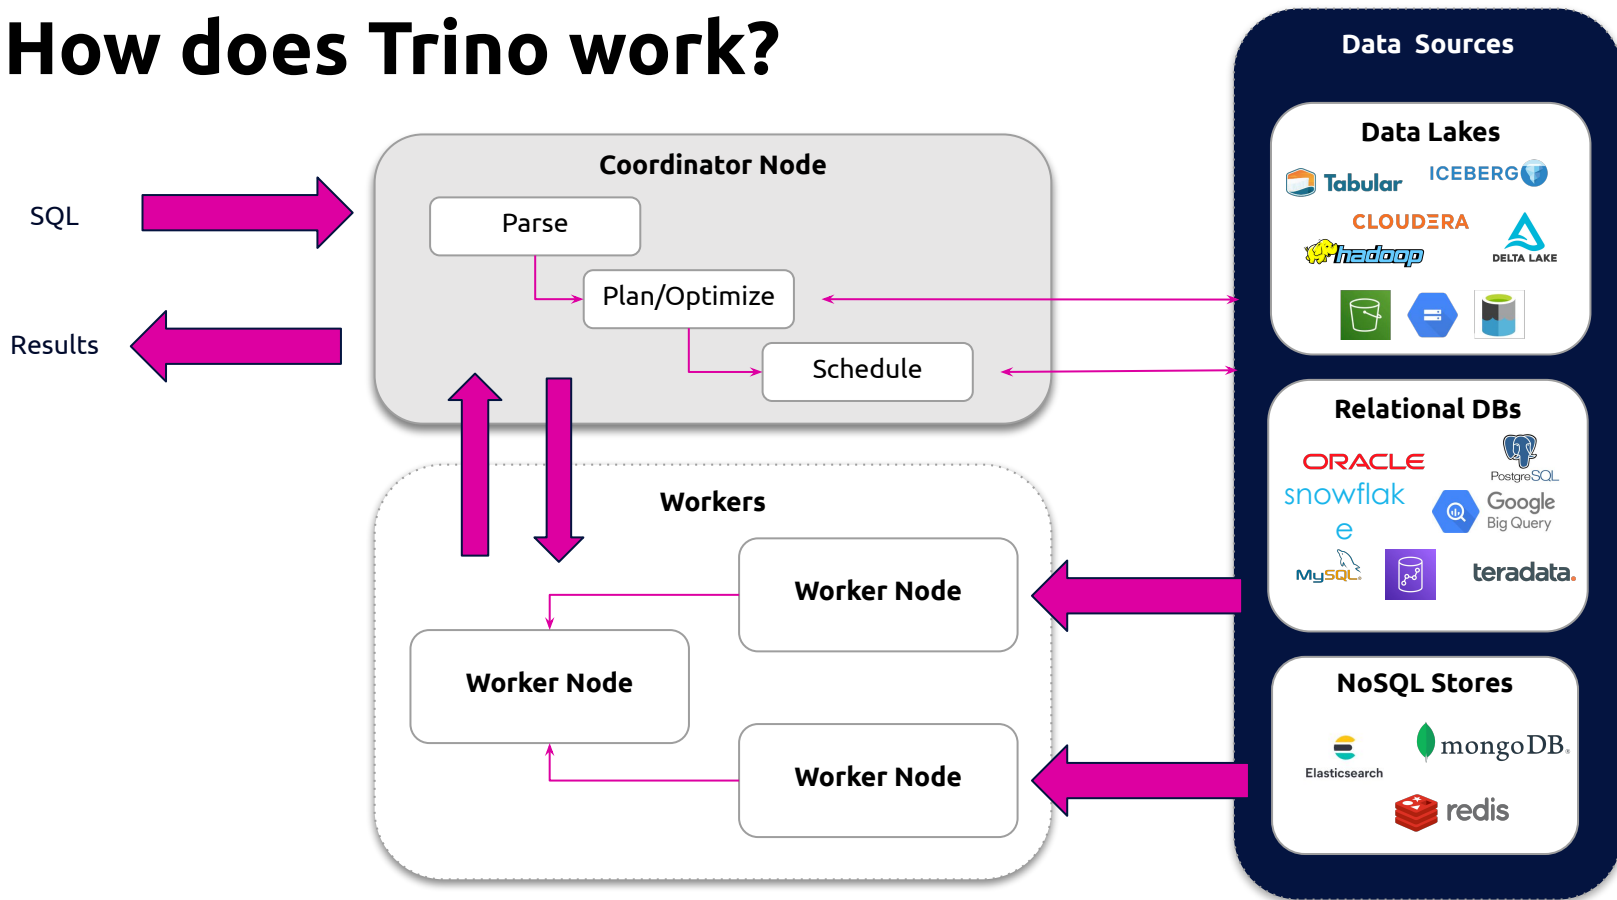
Trino (*formerly known as Presto*) is a fast distributed SQL query engine designed to query large data sets distributed over one or more heterogeneous data sources.

- Harnesses the power of distributed computing
- Separates compute from storage
- ANSI SQL compliant



<https://trino.io>

# How does Trino work?



<https://trino.io>




# Picking your components

Iceberg is the industry standard table format

# The Challenges of the invisible Hive “spec”

Hive has been critical for the evolution of SQL querying in distributed systems


- Rigid partitions – *requires user involvement*
- DIY schema evolution
- Not optimized for object storage – *need to scan all files in a “folder”*
- Bolt-on ACID transactions have always been squirrely – *inconsistency, correctness issue*
- Performance & scalability concerns with the metastore
- No inherent table content versioning



# Apache Iceberg

- Created by Ryan Blue & Daniel Weeks at Netflix in 2017
- Solve the challenges of performance, data modification and schema evolution in the lake + offer benefits of versioning
- Uses open data concepts (orc, parquet, avro) and architecture

## Multi-Engine Platform



# Iceberg: lake choice + warehouse behavior

## SQL behavior

- Schema and partition evolution
- Hidden partitioning

## Modern warehouse SQL

- UPDATE / DELETE / MERGE
- Time travel & rollback (via versioning)





# Iceberg: lake choice + warehouse behavior

## SQL behavior

- Schema and partition evolution
- Hidden partitioning

## Modern warehouse SQL

- UPDATE / DELETE / MERGE
- Time travel & rollback (via versioning)




# Architecture overview



snapshots are created anytime data or structure is changed


Metadata layer  
lives in the  
./metadata dir



Data files



Data files



Data files

saved in the  
./data dir  
Data layer

metadata and data persisted together on the data lake

# Iceberg should be invisible

## Avoid unpleasant surprises

- No zombie data
- Performance is not mysterious
- Reduced metastore reliance

## Doesn't steal attention

- Fast metadata operations
- Automate the boring stuff
- Fix problems without migration

## Optimistic Concurrency

- Allows multiple writes simultaneously, checks for conflicts before final commit

## Universal open standard



# Building a data lakehouse


# Open Data Lakehouse Benefits

## Data Warehouse Benefits



- ACID transactions
- Fined grained access control
- Data quality
- High performance and concurrency
- Highly curated data
- *Typically proprietary systems*
- Best for business intelligence use cases

## Data Lake Benefits



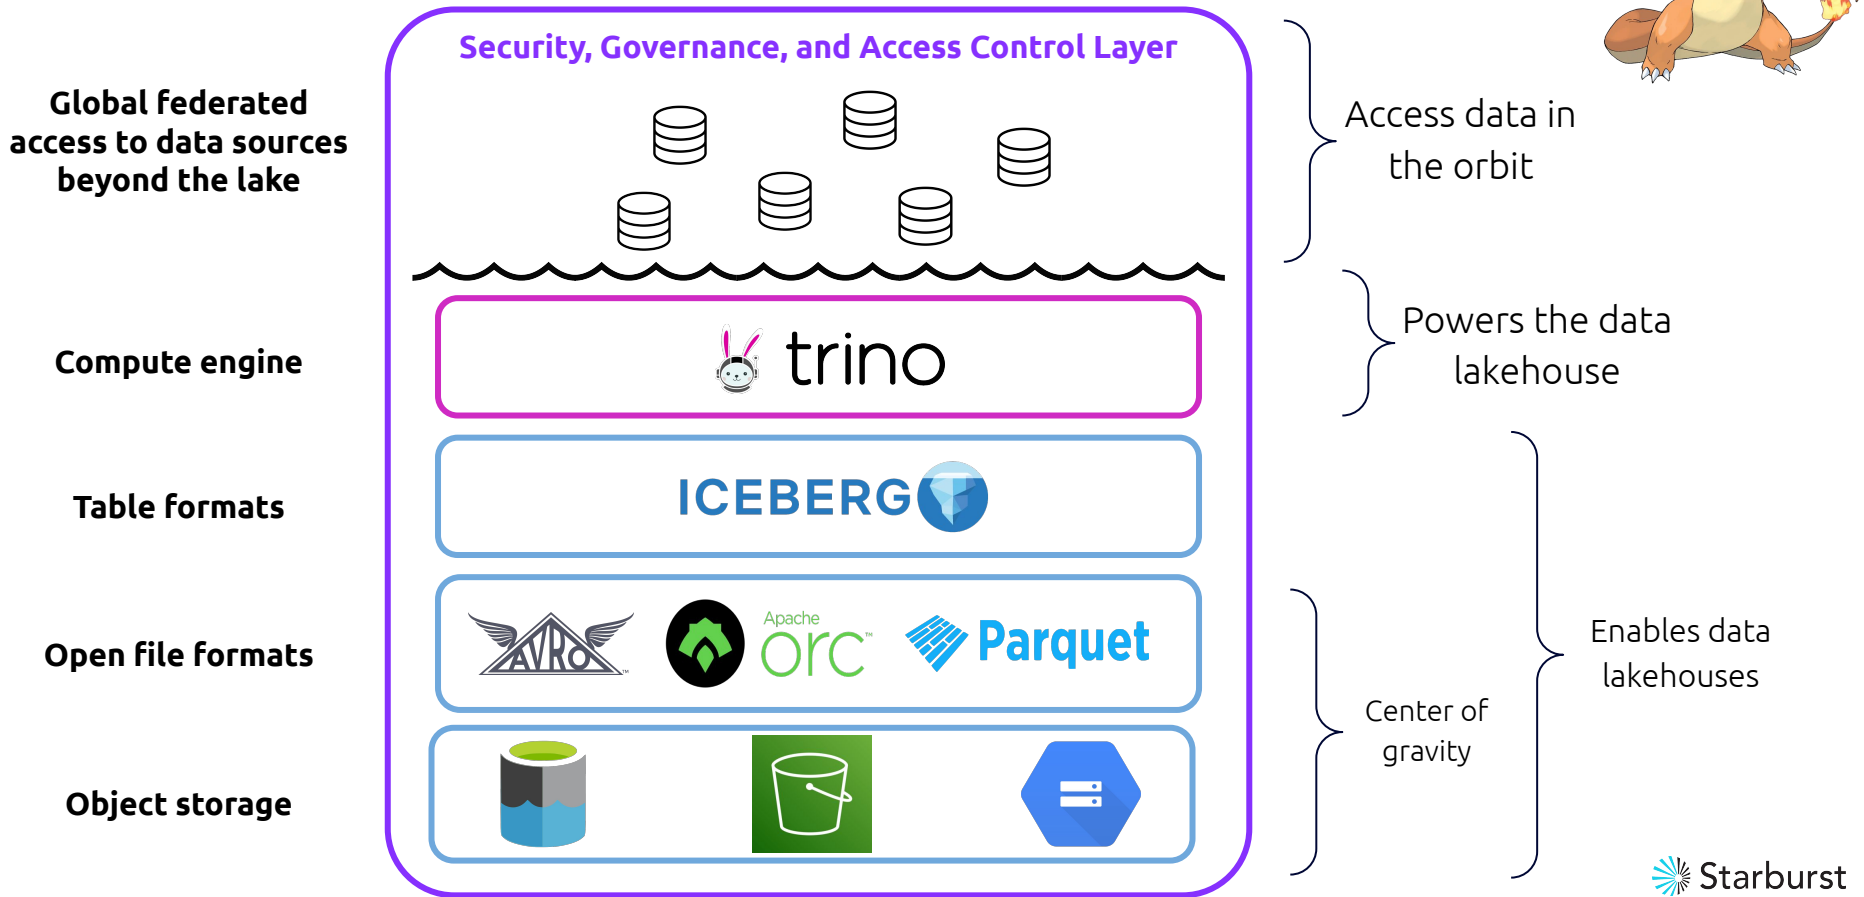
- Petabyte scale
- Cost efficient
- Open formats
- Separation of storage & compute
- Structured and unstructured data
- Best for data science and data engineering use cases



**Lakehouse = the doodle of data architecture**

*Apply data warehouse principles to the data lake of your choice*

# The Open Data Lakehouse



# When NOT to migrate from Hive

# Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*



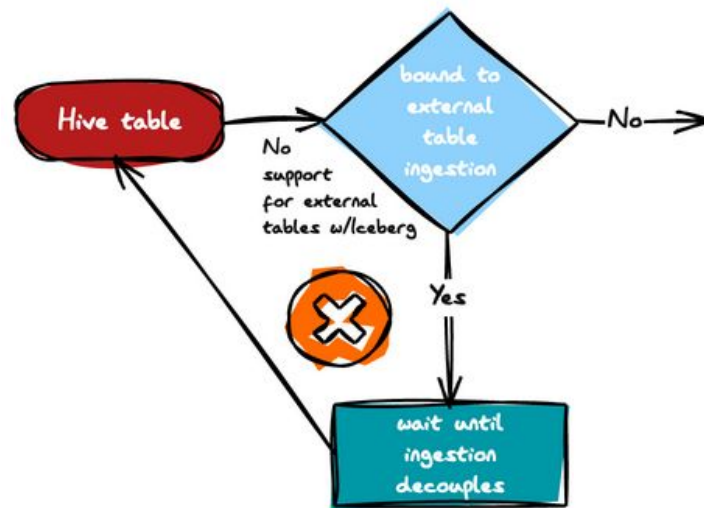


# Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

- **Iceberg requires tables to be managed** - no externally ingested files




# Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

- Iceberg requires tables to be managed - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution




# Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

- **Iceberg requires tables to be managed** - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- **High degree of confidence in existing performance and scalability** - testing validates no foreseeable concerns in the short to long term



# Reasons to NOT migrate



**Time and effort will be required - *make sure the juice is worth the squeeze***

- **Iceberg requires tables to be managed** - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- **High degree of confidence in existing performance and scalability** - testing validates no foreseeable concerns in the short to long term
- **Simply no resources available to migrate** - much less the necessary testing to find any unforeseen issues

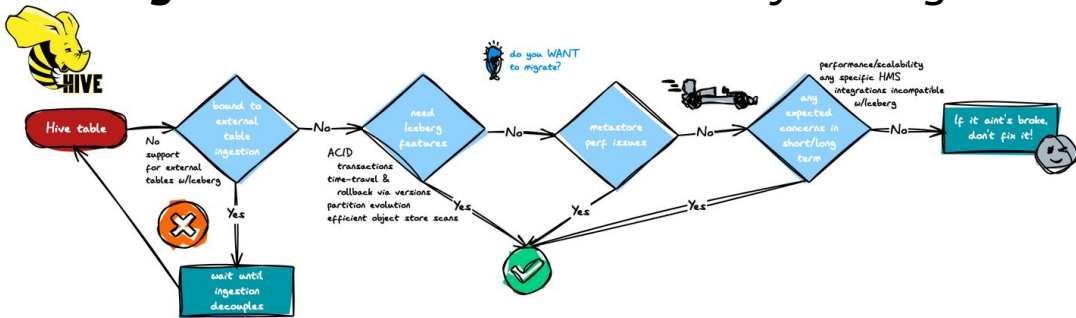


# Reasons to NOT migrate



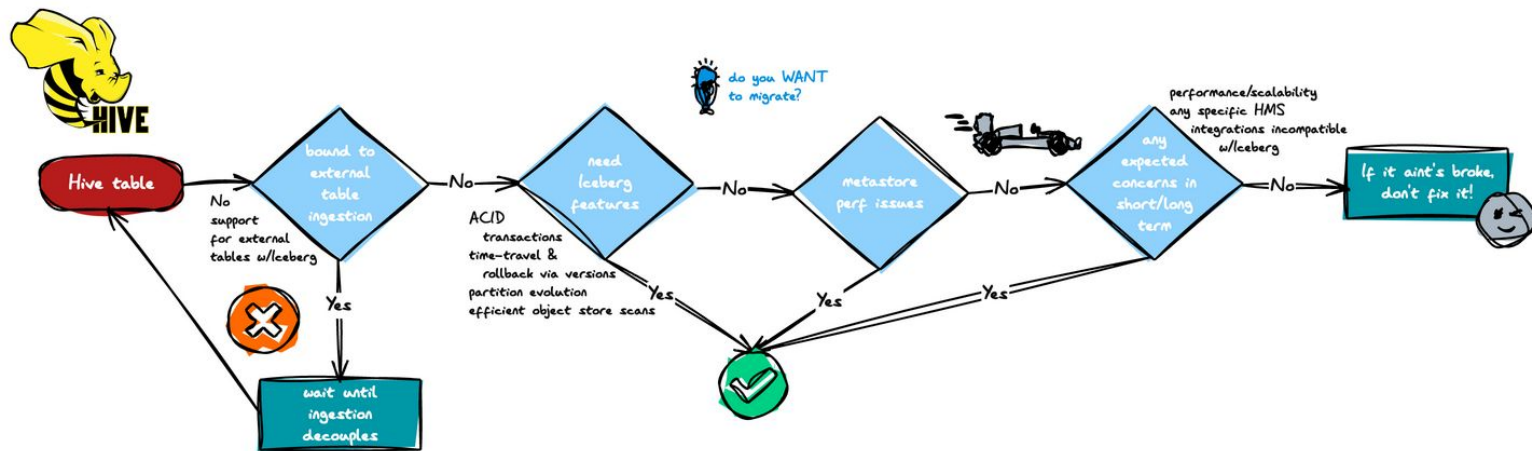
Time and effort will be required - *make sure the juice is worth the squeeze*

- **Iceberg requires tables to be managed** - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- **High degree of confidence in existing performance and scalability** - testing validates no foreseeable concerns in the short to long term
- **Simply no resources available to migrate** - much less the necessary testing to find any unforeseen issues




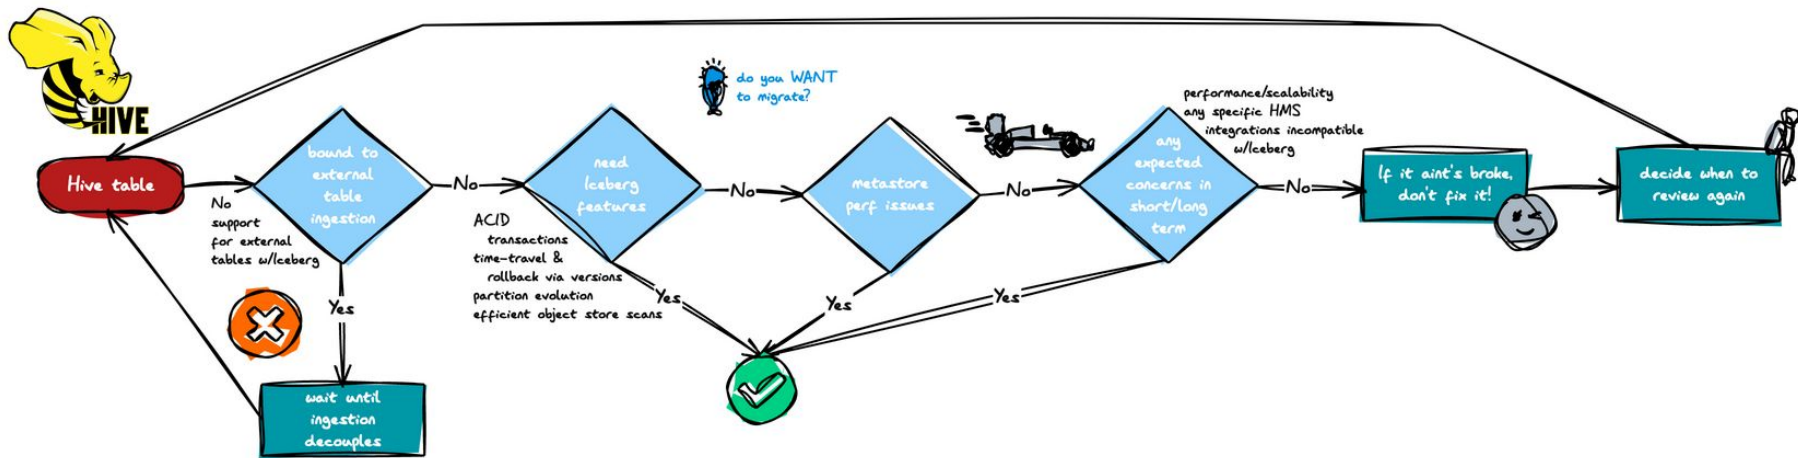
# Reasons to NOT migrate

Time and effort will be required - *make sure the juice is worth the squeeze*



# Reasons to NOT migrate (or not yet)

Time and effort will be required - *make sure the juice is worth the squeeze*



# Migration strategies

In-place vs shadow options



# Migration strategies

*Two approaches - let's define them*

## Shadow migration process

Creates a new Iceberg table modeled after the original Hive table whose values are then inserted into the new table; the original table can then be dropped

## The in-place method

Avoids rewriting the data files by modifying the table format type in the catalog and only building additional Iceberg metadata files

# In-place migration requirements


Takes over table's folder & creates Iceberg metadata - *recycles files*



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and  
TIMESTAMP is not compatible (millisecond vs microsecond precision)



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **bucketed\_by/bucket\_count are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **bucketed\_by/bucket\_count are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files





# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: `CHAR` and `TINYINT` not supported and `TIMESTAMP` is not compatible (millisecond vs microsecond precision)
- **`bucketed_by/bucket_count` are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: `CHAR` and `TINYINT` not supported and `TIMESTAMP` is not compatible (millisecond vs microsecond precision)
- **`bucketed_by/bucket_count` are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



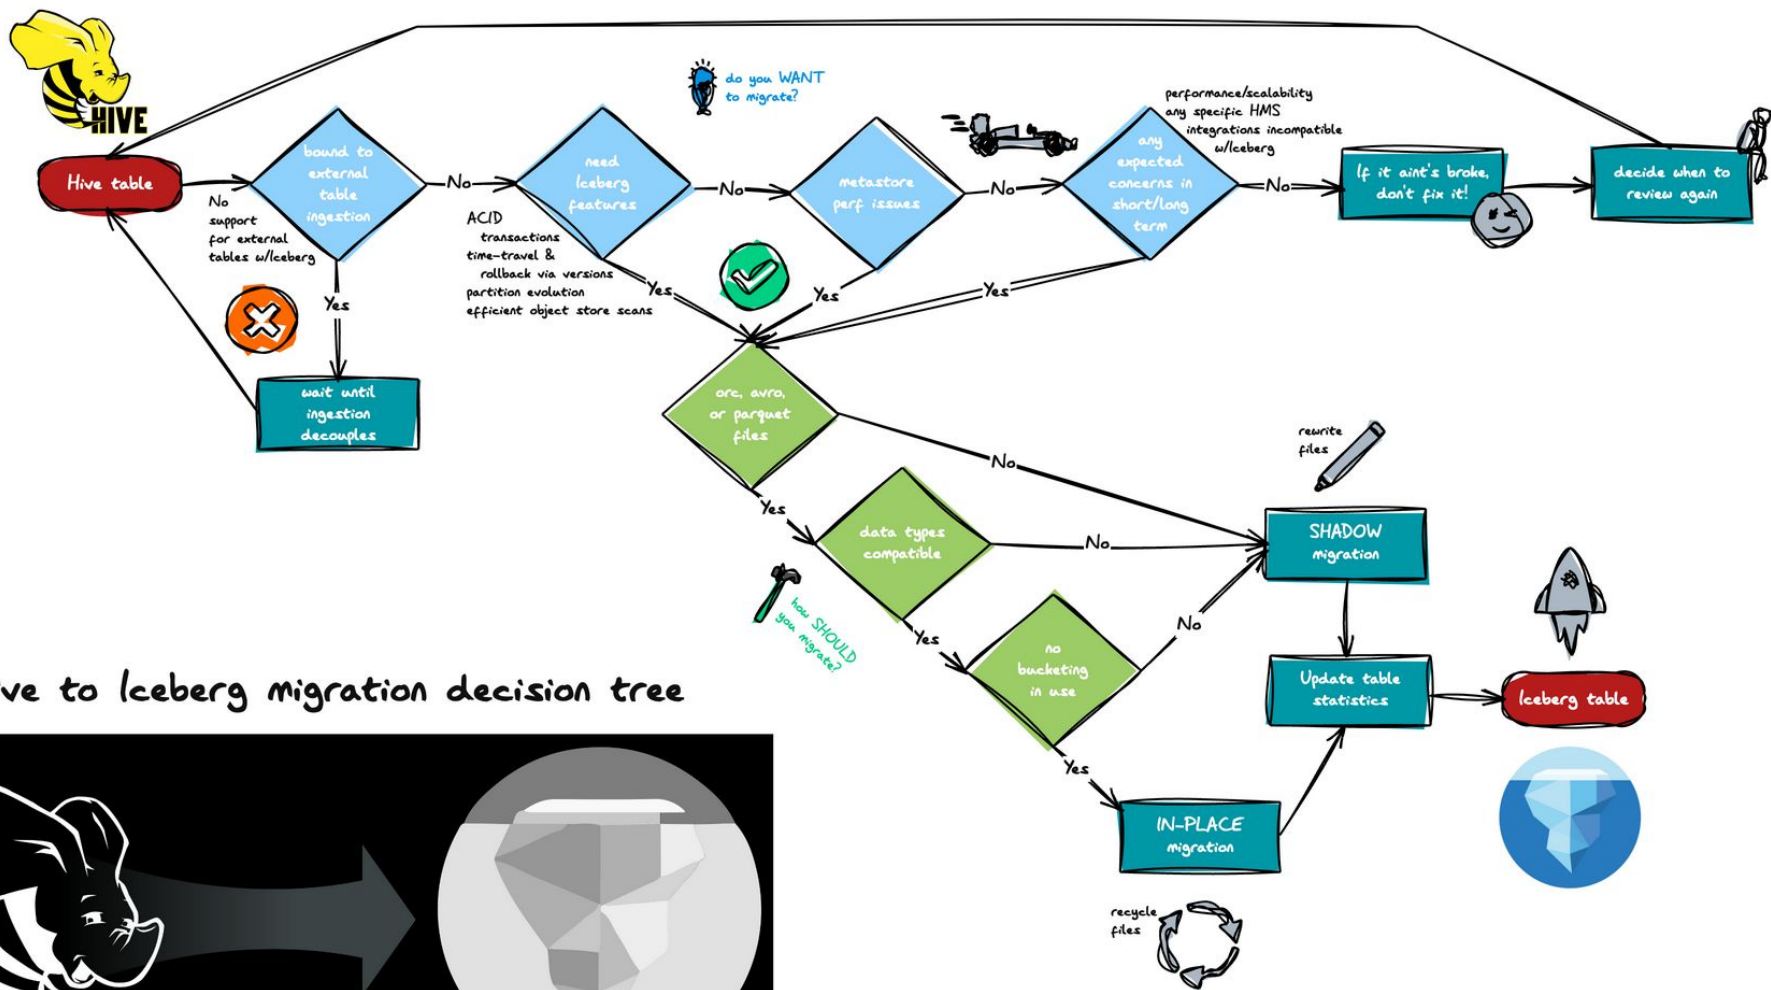
# In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*


- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: `CHAR` and `TINYINT` not supported and `TIMESTAMP` is not compatible (millisecond vs microsecond precision)
- **`bucketed_by/bucket_count` are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



**Put it all together**



Hive to Iceberg migration decision tree



# Additional considerations

# Migration considerations

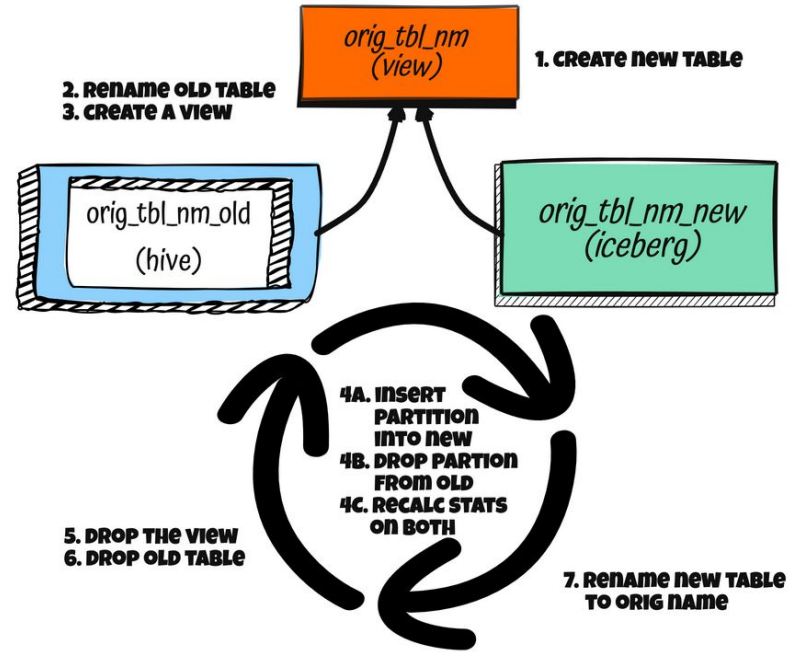
Any migration is an major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables

# Migration considerations

Any migration is an major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables





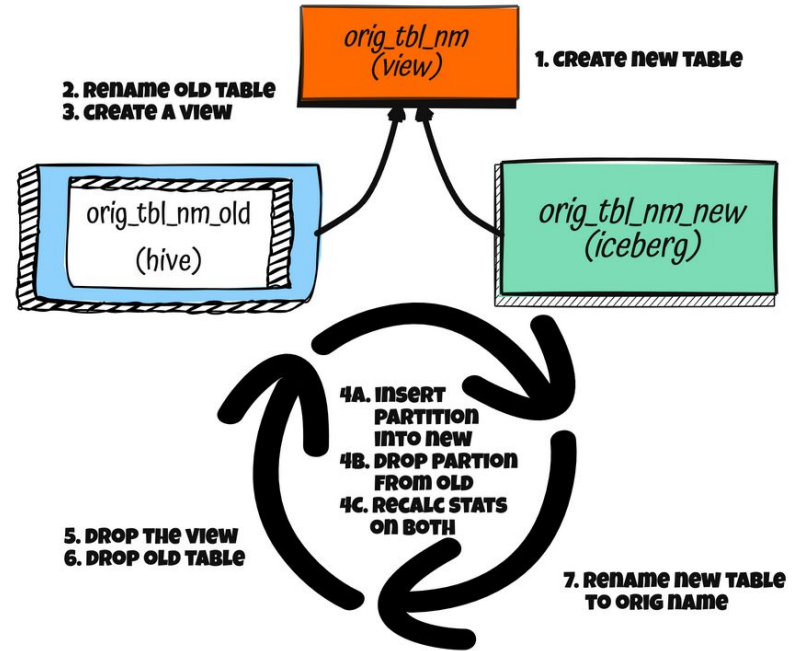
# Migration considerations

Any migration is an major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables



If time  
permits



# Next steps

# What are my next steps?

Decide if you are ready to begin & then engage Starburst for help

- Evaluate your existing data lake tables
- Consider tactical focus on largest tables vs comprehensive migration of all
- Visit <https://www.starburst.io/solutions/data-migrations/hive-iceberg/> for more information
- Get free guidance on your Hive to Iceberg migration by providing contact info at <https://www.starburst.io/info/hive-to-iceberg-migration-guidance/>

Demo artifacts at [https://github.com/lestermartin/events/tree/main/2024-05-08\\_Hive2Iceberg](https://github.com/lestermartin/events/tree/main/2024-05-08_Hive2Iceberg)



**Thank you!**

