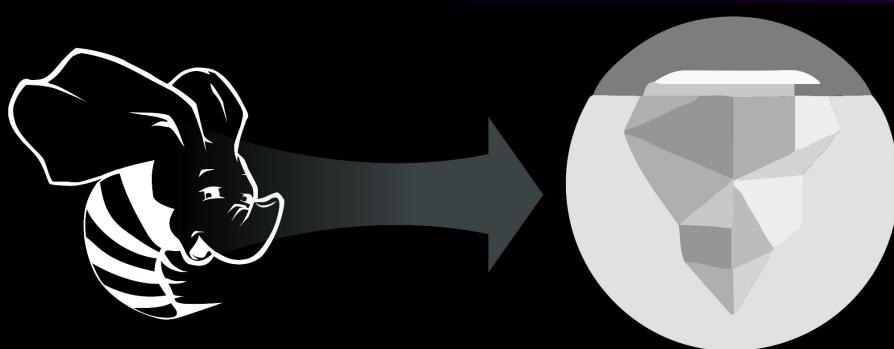




Hive to Iceberg

To Migrate, or Not to Migrate

Starburst Webinar



A portrait of Lester Martin, a middle-aged man with curly grey hair and glasses, smiling. He is wearing a dark t-shirt. Behind him is a circular graphic featuring stylized yellow and red shapes, with the word "SIMS" at the top and "LESTER KASAI" in large letters below it.

Connection before content

Lester Martin – <https://lestermartin.blog/about>

- Developer advocate @ Starburst
 - Blogging & forums
 - Webinars & videos
 - User groups & events
 - Training & tutorials
- 30+ years of technology experience
 - Started journey on TRS-80 Model III
 - Played most roles, but a programmer at my core
 - ½ career in OLTP and ½ in data analytics
 - Decade+ of “big data” experience to include
 - Trino/Starburst, Hadoop, Hive, Spark
 - NiFi, Kafka, Storm, Flink
 - HBase, MongoDB

devrel@starburst.io

Webinar Agenda

Slides, but DEMOs, too!

- Evolution of a data lakehouse
(the 3 min version)
- Picking your components
- Building a data lakehouse
- When NOT to migrate from Hive
- Migration strategies
- Additional considerations

Scan for a Trino and Iceberg cheat sheet



Evolution of the data lakehouse

How did we get here?

Data Architecture Evolution

Data Warehouse



Charmander

Data Lake



Charmeleon

Data Lakehouse



Charizard

The Data Warehouse



Popularized in the 90's to provide a 360 degree view

The Good

- Integrates siloed RDBMS's into one "centralized" location
- Simple & reliable analytical querying
- Data audit, governance and lineage
- Great for small amounts of data

The Bad

- Inability to store unstructured data
- Lack scalability and flexibility
- Tightly coupled storage and compute
- Expensive, proprietary hardware and software (*creating vendor lock-in*)

The Data Lake



Born out of the internet age and big data boom

The Good

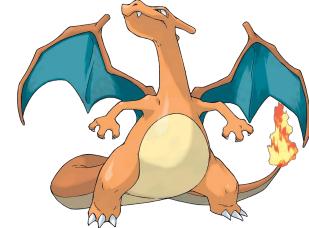
- In 2006, Apache Hadoop emerges so unstructured data can be processed at a scale previously unimaginable
- Shift toward parallel processing
- Capitalize on low cost object storage
- Allows for greater flexibility (schema on read)

The Bad

- Inability to support transactions, updates, or modifications
- Difficult to get top-tier performance
- Lack of data quality and inconsistent data formats
- Insufficient data lineage and limited data discoverability

The Data Lakehouse

Applying data warehouse principles to the data lake



- Utilize the ***separation of storage and compute*** to apply the reliability, performance, and data quality of the data warehouse to the openness and scalability of the data lake
- ***Increased performance and scalability*** through the use of indexing and caching via your query engine (**Trino**) and modern table formats (ex: **Iceberg**)
- Provide traditional ***data modifications*** (ex: UPDATE & MERGE commands) with ***ACID transaction*** guarantees over files stored in the data lake
- Tackle ***unstructured, semi-structured, and structured*** analytical data all in a data lakehouse - creating a place for AI/ML & BI use cases alike

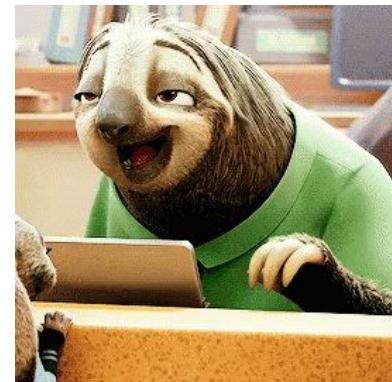


Picking your components

Trino is the best query engine ever

The data accessibility problem

Data practitioners faced the same challenges at Facebook in 2010



- Facebook created Hive to query terabytes of data in Hadoop using SQL
- Data scientists attempted to query massive object stores, but performance was too slow
- Data consumers were limited by the number of queries they could run — often ***fewer than 10*** in one day

Enter Trino (Presto)

A new open source query engine designed for speed

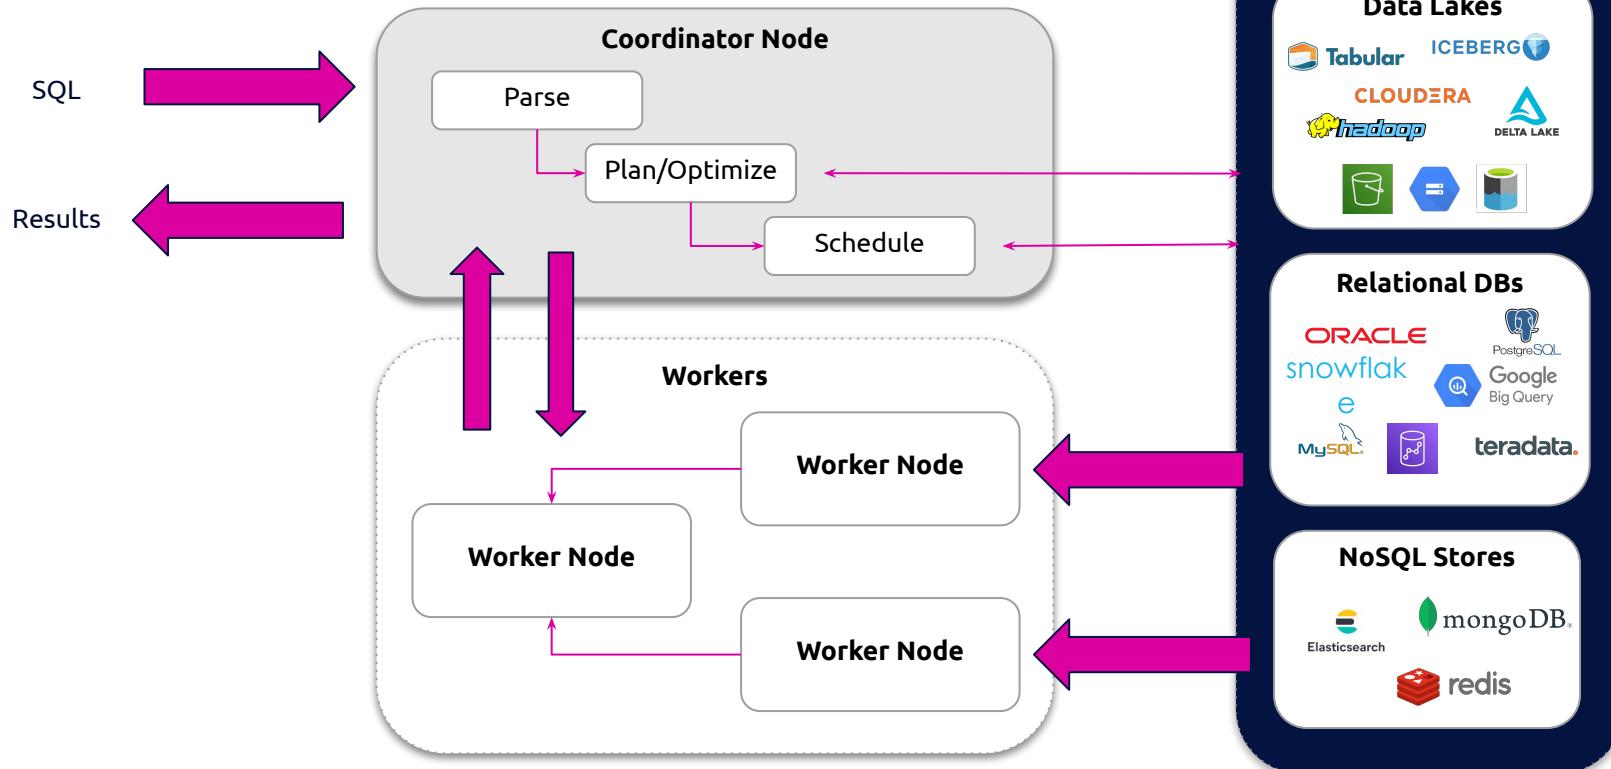
Trino (*formerly known as Presto*) is a fast distributed SQL query engine designed to query large data sets distributed over one or more heterogeneous data sources.

- Harnesses the power of distributed computing
- Separates compute from storage
- ANSI SQL compliant



<https://trino.io>

How does Trino work?



<https://trino.io>

 Starburst



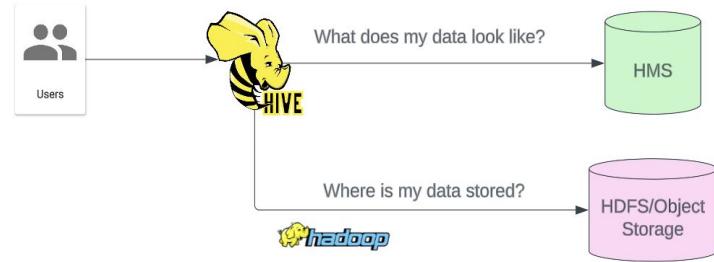
Picking your components

Iceberg is the industry standard table format

The Challenges of the invisible Hive “spec”

Hive has been critical for the evolution of SQL querying in distributed systems

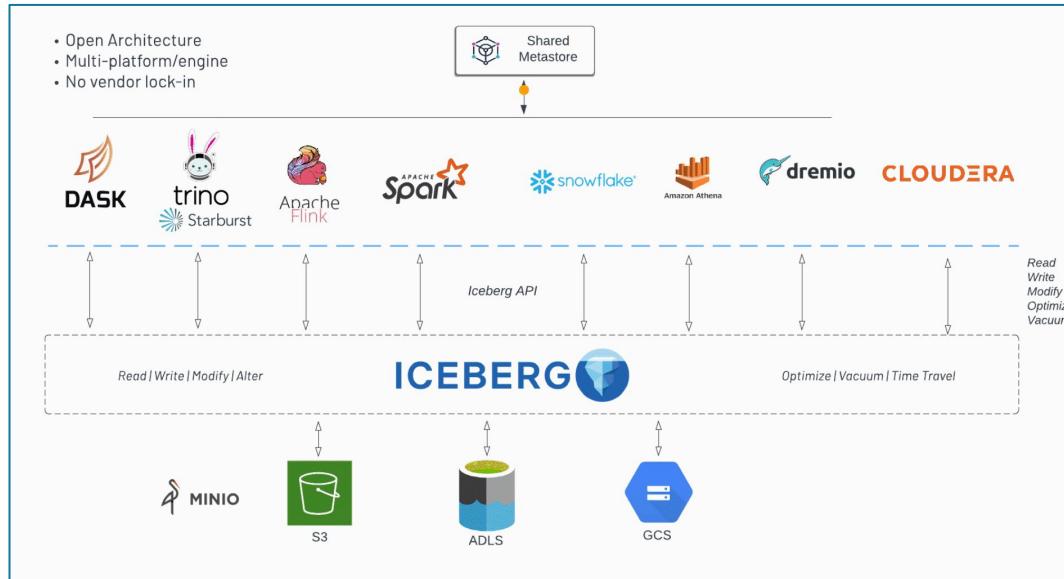
- Rigid partitions – *requires user involvement*
- DIY schema evolution
- Not optimized for object storage – *need to scan all files in a “folder”*
- Bolt-on ACID transactions have always been squirrelly – *inconsistency, correctness issue*
- Performance & scalability concerns with the metastore
- No inherent table content versioning



Apache Iceberg

- Created by Ryan Blue & Daniel Weeks at Netflix in 2017
- Solve the challenges of performance, data modification and schema evolution in the lake + offer benefits of versioning
- Uses open data concepts (orc, parquet, avro) and architecture

Multi-Engine Platform



Iceberg: lake choice + warehouse behavior

SQL behavior

- Schema and partition evolution
- Hidden partitioning

Modern warehouse SQL

- UPDATE / DELETE / MERGE
- Time travel & rollback (via versioning)



Iceberg: lake choice + warehouse behavior

SQL behavior

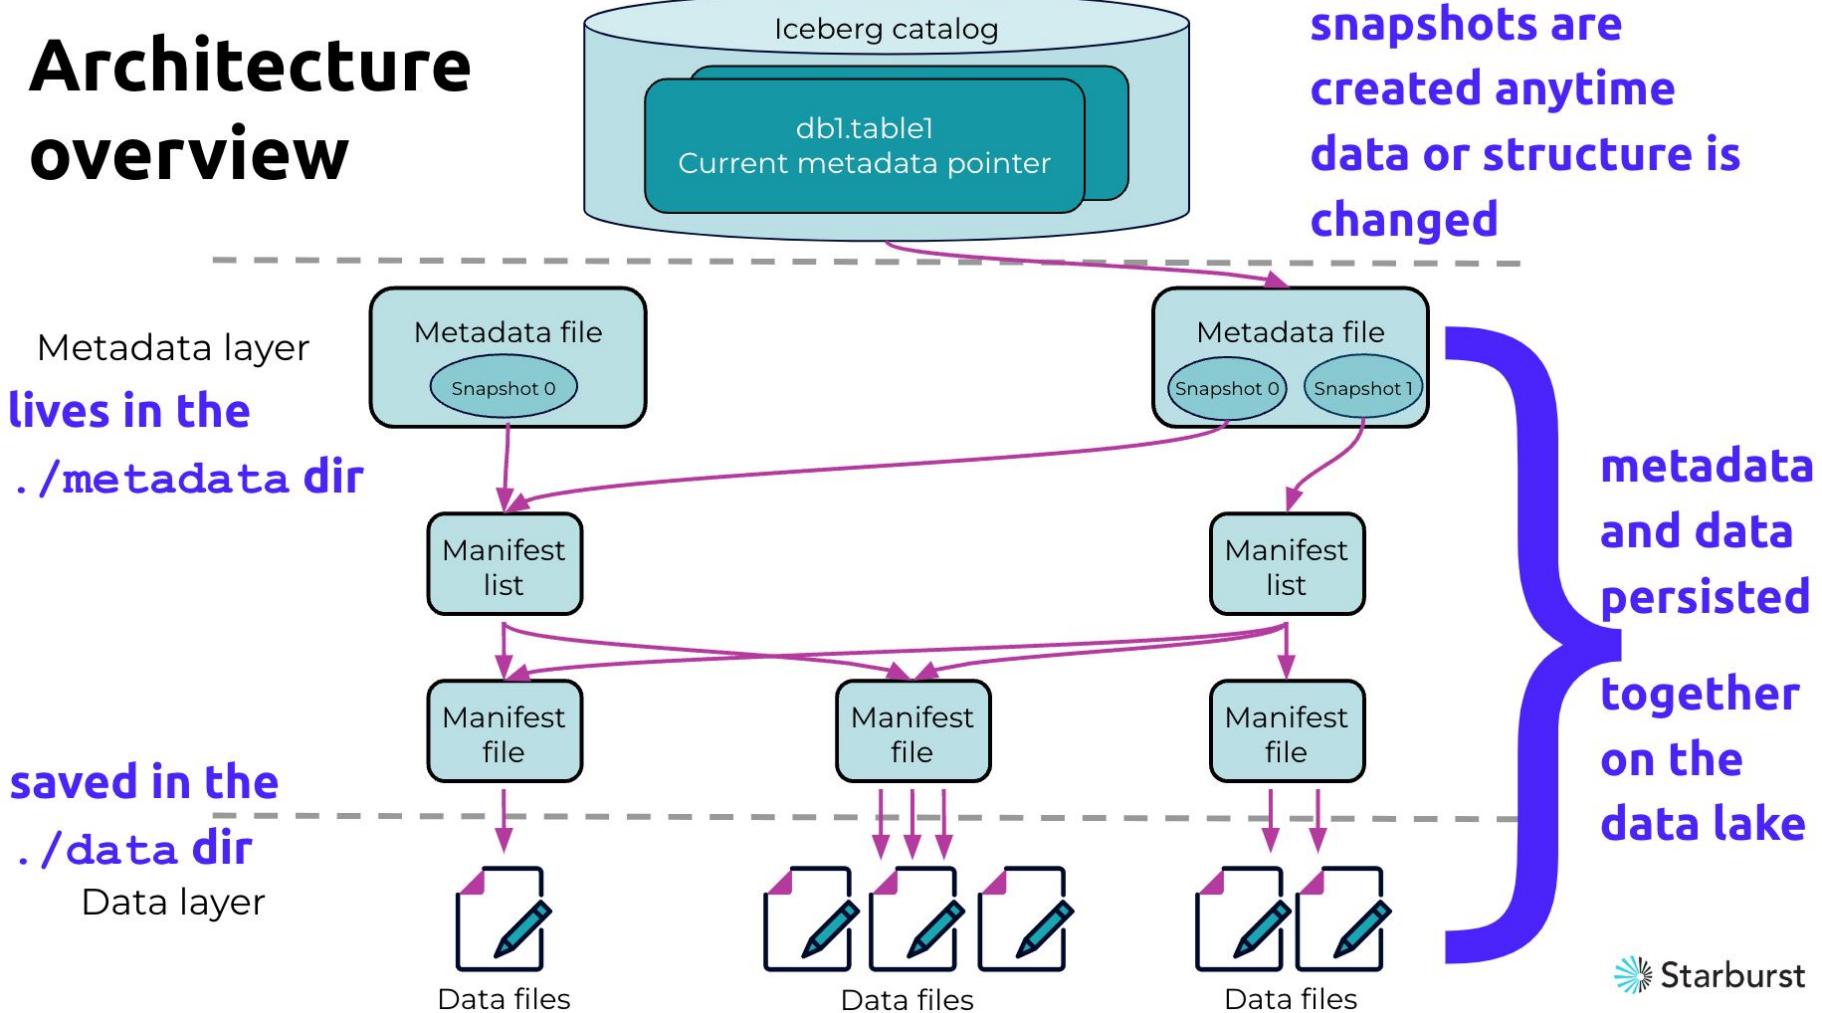
- Schema and partition evolution
- Hidden partitioning

Modern warehouse SQL

- UPDATE / DELETE / MERGE
- Time travel & rollback (via versioning)



Architecture overview



Iceberg should be invisible

Avoid unpleasant surprises

- No zombie data
- Performance is not mysterious
- Reduced metastore reliance

Doesn't steal attention

- Fast metadata operations
- Automate the boring stuff
- Fix problems without migration

Optimistic Concurrency

- Allows multiple writes simultaneously, checks for conflicts before final commit

Universal open standard



Building a data lakehouse

Open Data Lakehouse Benefits

Data Warehouse Benefits

- ACID transactions
- Fined grained access control
- Data quality
- High performance and concurrency
- Highly curated data
- *Typically proprietary systems*
- Best for business intelligence use cases



Data Lake Benefits

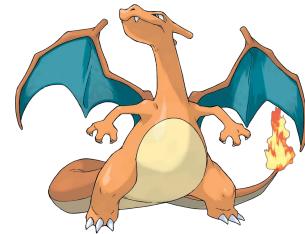
- Petabyte scale
- Cost efficient
- Open formats
- Separation of storage & compute
- Structured and unstructured data
- Best for data science and data engineering use cases



Lakehouse = the doodle of data architecture

Apply data warehouse principles to the data lake of your choice

The Open Data Lakehouse



Global federated access to data sources beyond the lake

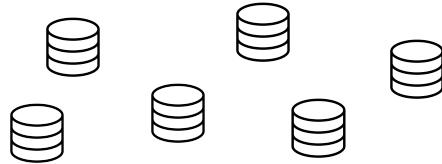
Compute engine

Table formats

Open file formats

Object storage

Security, Governance, and Access Control Layer



Access data in the orbit

Powers the data lakehouse

Enables data lakehouses

Center of gravity

When NOT to migrate from Hive

Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

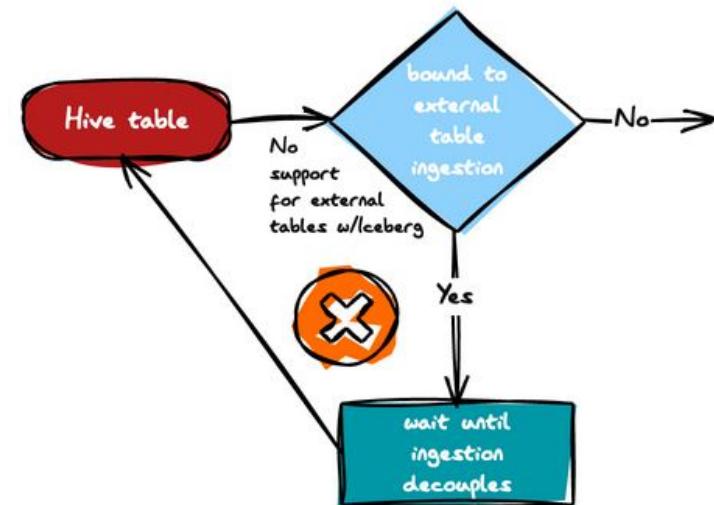


Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

- Iceberg requires tables to be managed - no classical external tables



Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

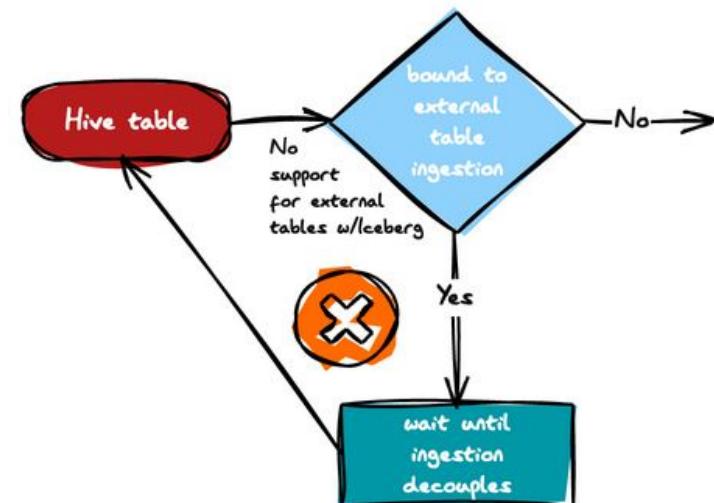
- **Iceberg requires tables to be managed** - no classical external tables

NOTE: DIY integration via add_files procedure

The add_files procedure supports adding files, and therefore the contained data, to a target table, specified after ALTER TABLE. It loads the files from a object storage path specified with the required location parameter. The files must use the specified format, with ORC and PARQUET as valid values. The target Iceberg table must use the same format as the added files. The procedure does not validate file schemas for compatibility with the target Iceberg table. The location property is supported for partitioned tables.

The following examples copy ORC-format files from the location s3://my-bucket/a/path into the Iceberg table iceberg_customer_orders in the lakehouse schema of the example catalog:

```
ALTER TABLE example.lakehouse.iceberg_customer_orders
EXECUTE add_files(
  location => 's3://my-bucket/a/path',
  format => 'ORC')
```

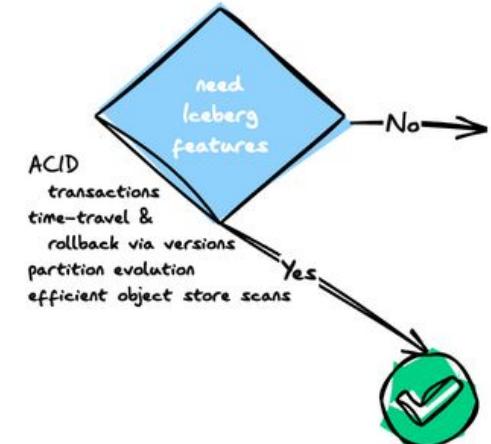


Reasons to NOT migrate



Time and effort will be required - make sure the juice is worth the squeeze

- Iceberg requires tables to be managed - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution

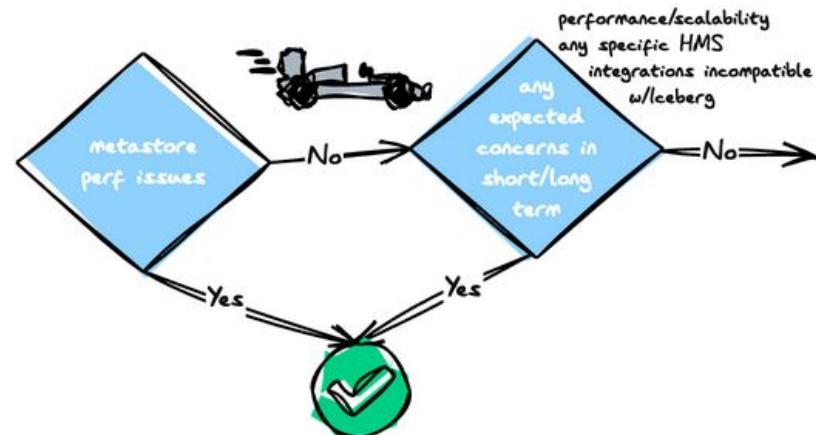


Reasons to NOT migrate



Time and effort will be required - *make sure the juice is worth the squeeze*

- Iceberg requires tables to be managed - no externally ingested files
- Functional benefits not valuable enough - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- **High degree of confidence in existing performance and scalability** - testing validates no foreseeable concerns in the short to long term



Reasons to NOT migrate



Time and effort will be required - make sure the juice is worth the squeeze

- Iceberg requires tables to be managed - no externally ingested files
- Functional benefits not valuable enough - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- High degree of confidence in existing performance and scalability - testing validates no foreseeable concerns in the short to long term
- **Simply no resources available to migrate** - much less the necessary testing to find any unforeseen issues

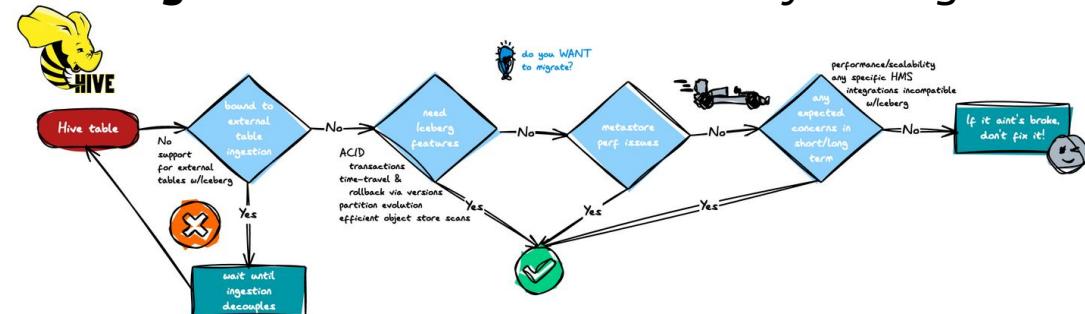


Reasons to NOT migrate



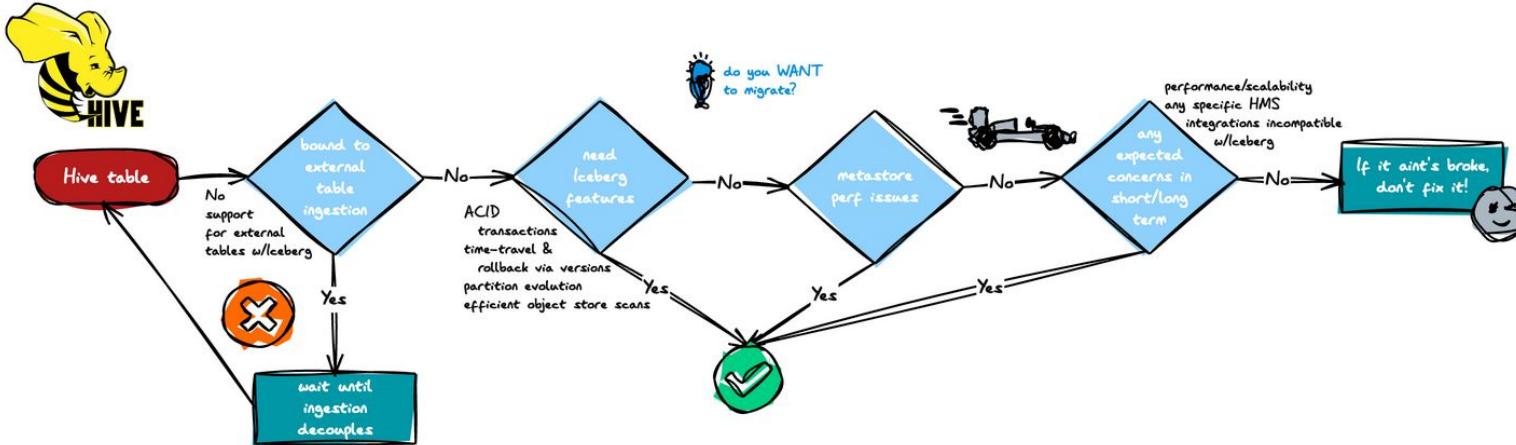
Time and effort will be required - *make sure the juice is worth the squeeze*

- **Iceberg requires tables to be managed** - no externally ingested files
- **Functional benefits not valuable enough** - ACID transactions, versioning, hidden partitioning, schema/partition evolution
- **High degree of confidence in existing performance and scalability** - testing validates no foreseeable concerns in the short to long term
- **Simply no resources available to migrate** - much less the necessary testing to find any unforeseen issues



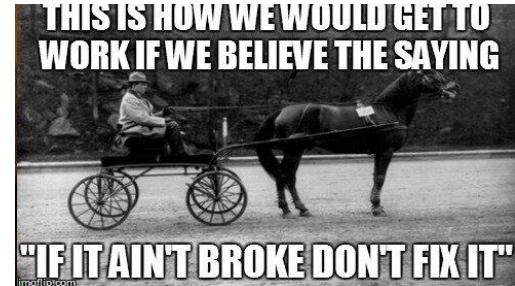
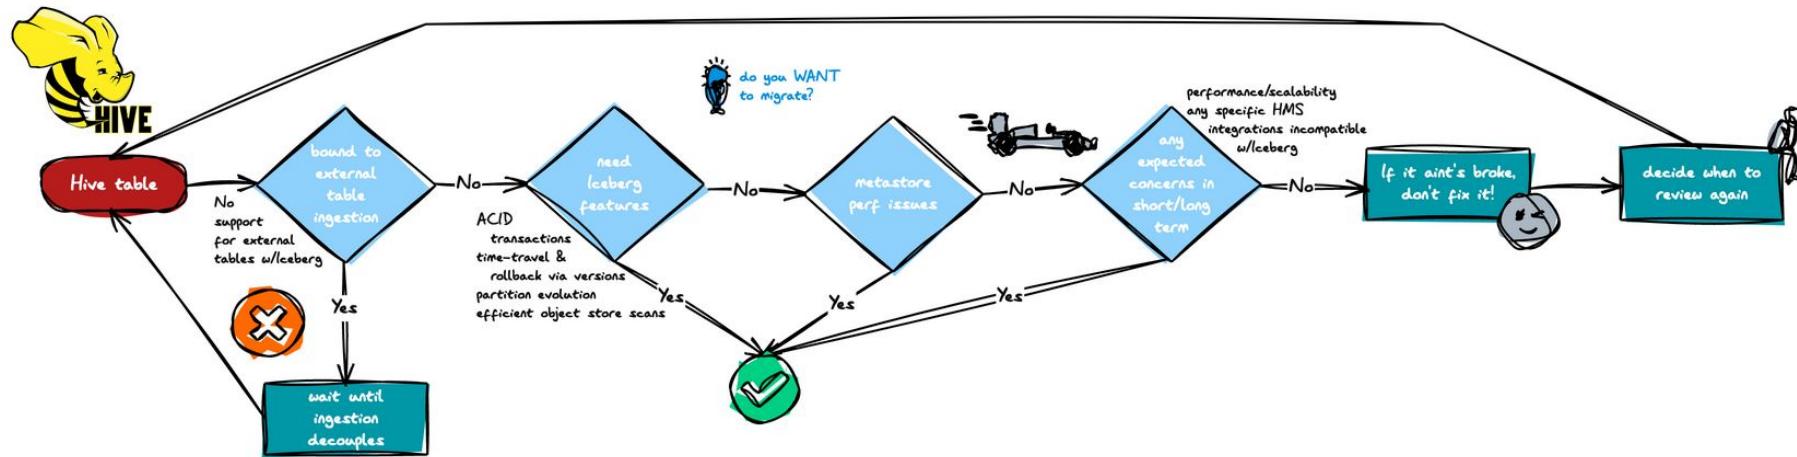
Reasons to NOT migrate

Time and effort will be required - *make sure the juice is worth the squeeze*



Reasons to NOT migrate (or not yet)

Time and effort will be required - *make sure the juice is worth the squeeze*



Migration strategies

In-place vs shadow options

Migration strategies

Two approaches - let's define them

Shadow migration process

Creates a new Iceberg table modeled after the original Hive table whose values are then inserted into the new table; the original table can then be dropped

The in-place method

Avoids rewriting the data files by modifying the table format type in the catalog and only building additional Iceberg metadata files

In-place migration requirements

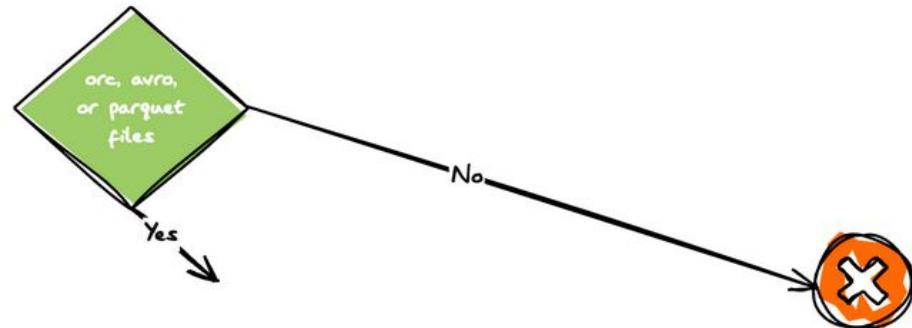
Takes over table's folder & creates Iceberg metadata - *recycles files*



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

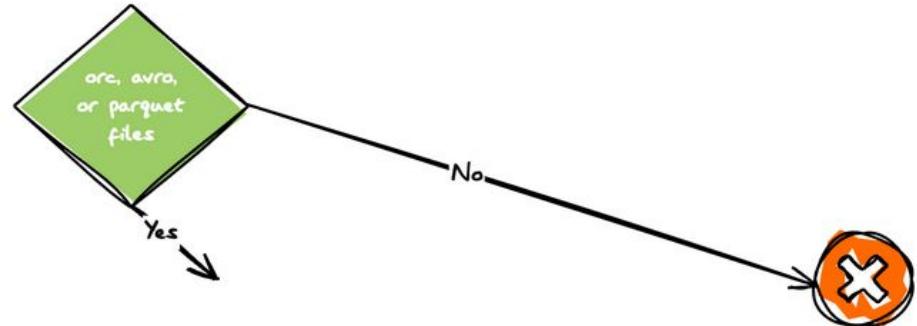
- **Compatible data file formats** - Parquet, ORC & Avro



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

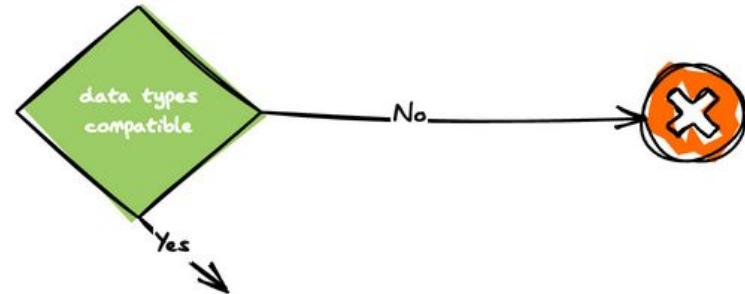
- Compatible data file formats - Parquet, ORC & Avro



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

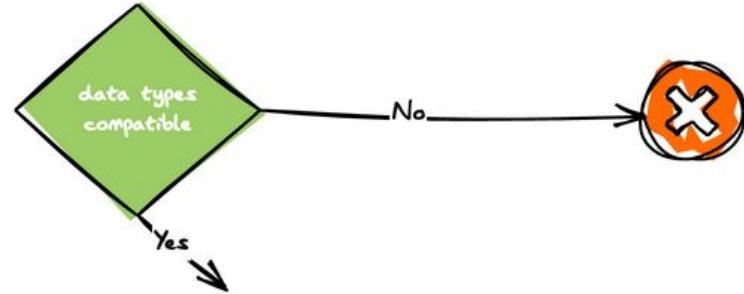
- Compatible data file formats - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

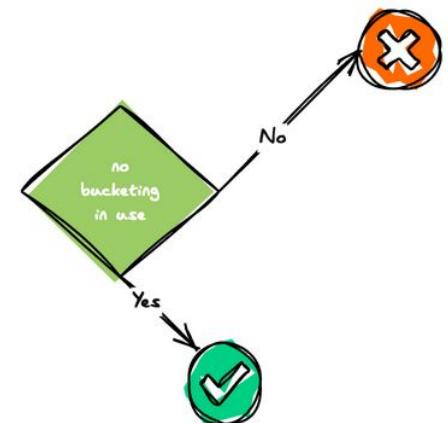
- Compatible data file formats - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

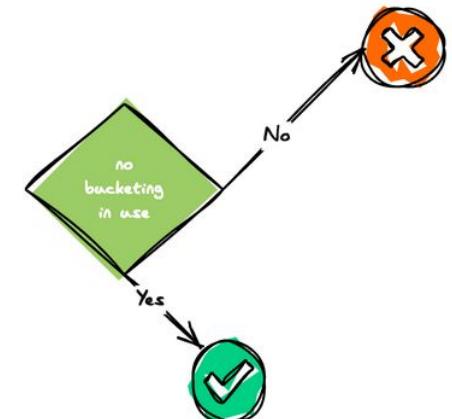
- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **`bucketed_by/bucket_count` are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **`bucketed_by/bucket_count` are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **bucketed_by/bucket_count are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **bucketed_by/bucket_count are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



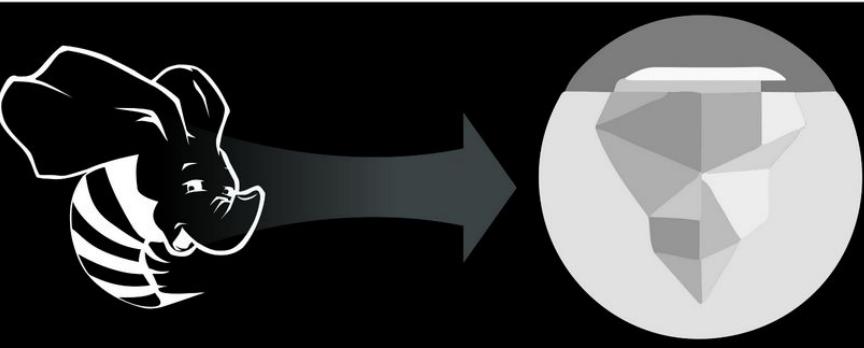
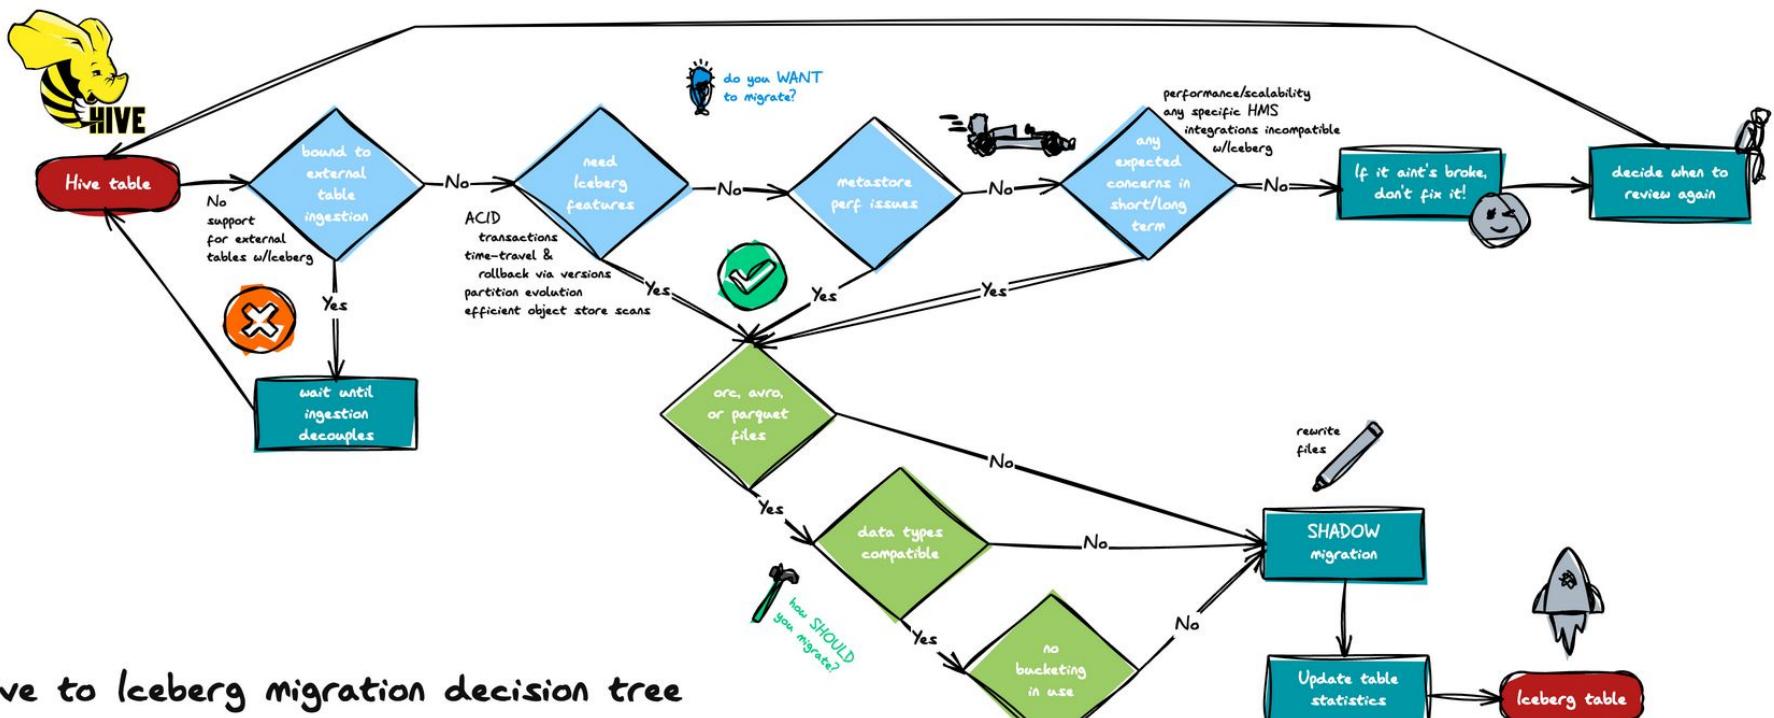
In-place migration requirements

Takes over table's folder & creates Iceberg metadata - *recycles files*

- **Compatible data file formats** - Parquet, ORC & Avro
- **Compatible column data types** - ex: CHAR and TINYINT not supported and TIMESTAMP is not compatible (millisecond vs microsecond precision)
- **bucketed_by/bucket_count are ignored** - Iceberg implements bucketing as numbered partitioned folders, not files



Put it all together



Additional considerations

Migration considerations

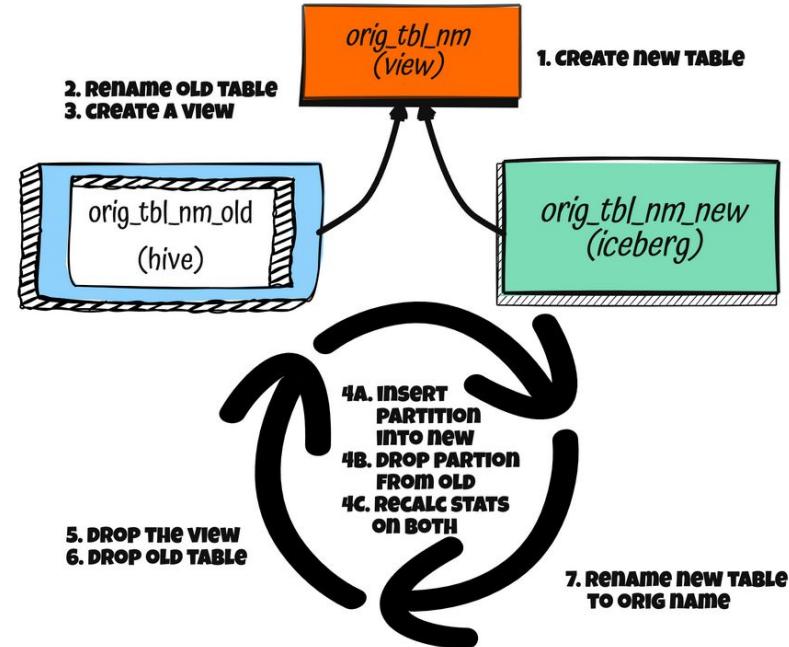
Any migration is a major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables

Migration considerations

Any migration is a major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables



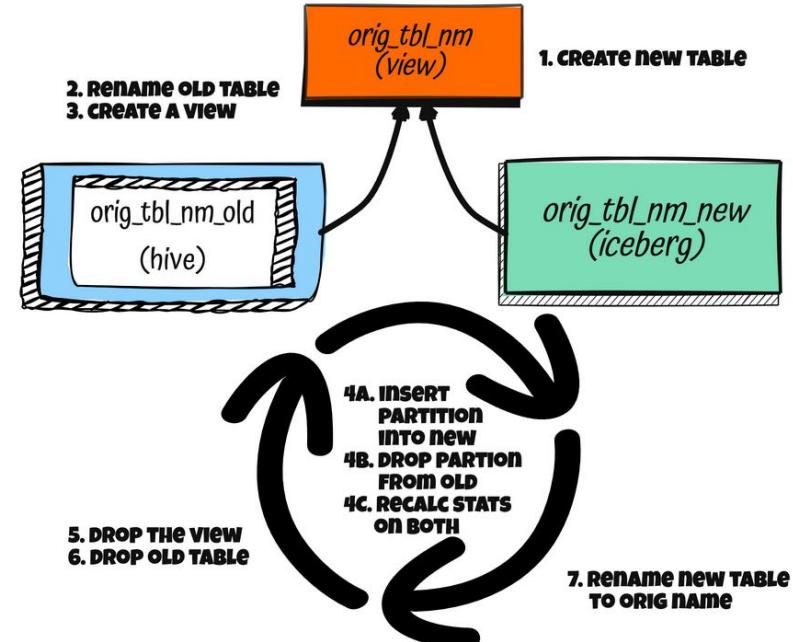
Migration considerations

Any migration is a major event and should be planned & tested accordingly

- Test, test, and more test - *don't forget about a backout strategy*
- Automate maintenance activities
- Consider staging rewrites for very large, heavily-partitioned, tables



If time
permits



Next steps

What are my next steps?

Decide if you are ready to begin & then engage Starburst for help

- Evaluate your existing data lake tables
- Consider tactical focus on largest tables vs comprehensive migration of all
- Visit <https://www.starburst.io/solutions/data-migrations/hive-iceberg/> for more information
- Get free guidance on your Hive to Iceberg migration by providing contact info at <https://www.starburst.io/info/hive-to-iceberg-migration-guidance/>

Demo artifacts at https://github.com/lestermartin/events/tree/main/2025-03-04_Hive2Iceberg



Thank you!



Starburst