

Apache Iceberg ingestion with Apache NiFi

Lester Martin; Trino Developer Advocate @ Starburst

Berlin Buzzwords
2025



Connection before content

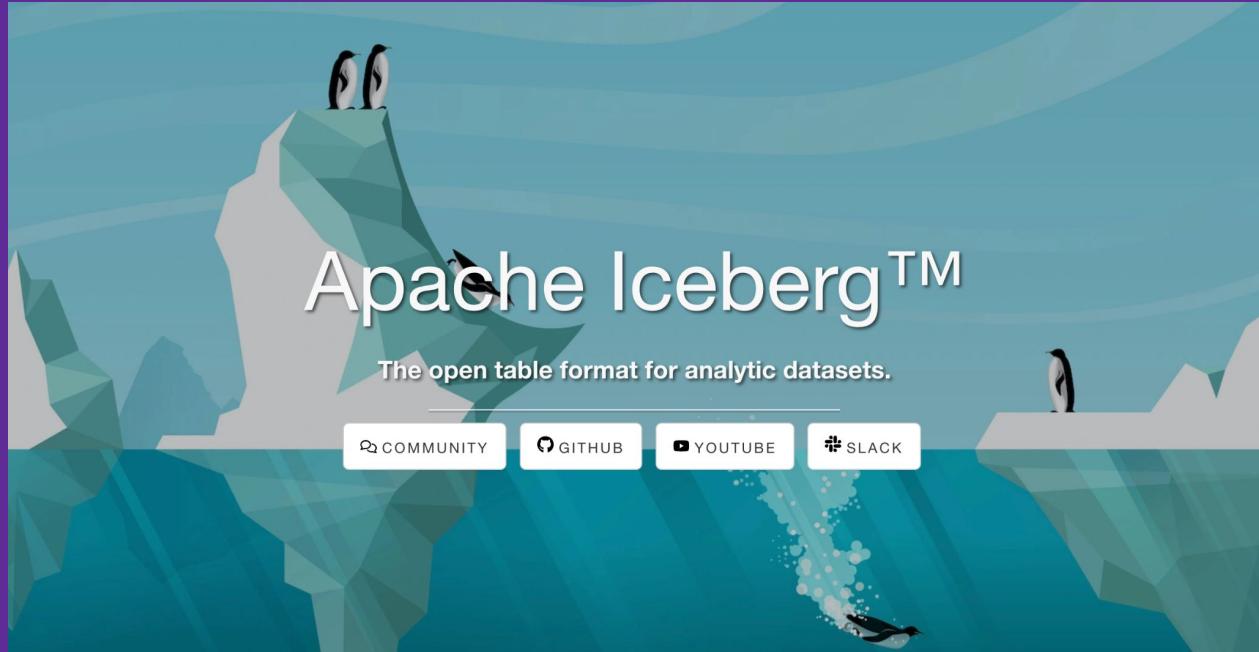
Lester Martin – <https://linktr.ee/lestermartin>

- Trino developer advocate @ Starburst
- 30+ years of technology experience
 - Started my journey on a TRS-80 Model III
 - Played most every role, but a programmer at my core
 - ½ career in OLTP and ½ in data analytics
 - Decade+ of “big data” experience to include
 - Trino/Starburst, Hadoop, Hive, Spark
 - NiFi, Kafka, Storm, Flink
 - HBase, MongoDB



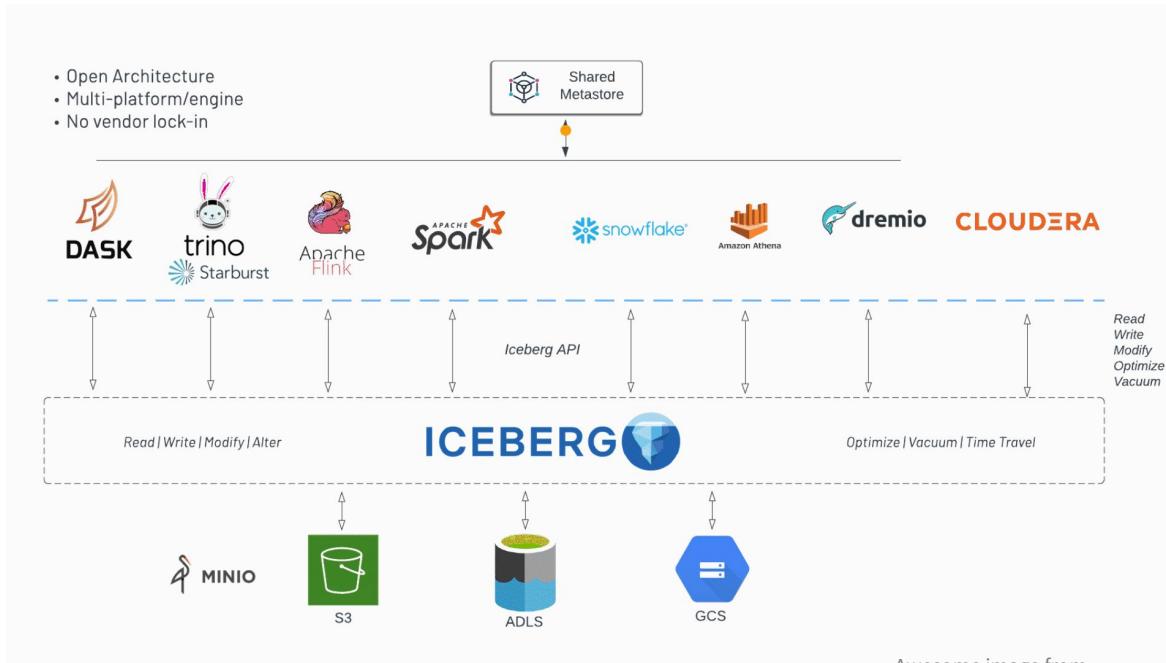
Apache Iceberg

What is it?



Multi-engine support for ludicrously large tables

Iceberg is a **high-performance** format for **huge analytic tables**. Iceberg brings reliability and simplicity of SQL tables to big data, while making it possible **for engines** like Spark, Trino, Presto, Hive and Impala **to safely work with the same tables, at the same time**.



ACID-compliant transactions

INSERT, UPDATE, DELETE, **and MERGE** statements

Only single-statement boundaries (aka auto-commit) today



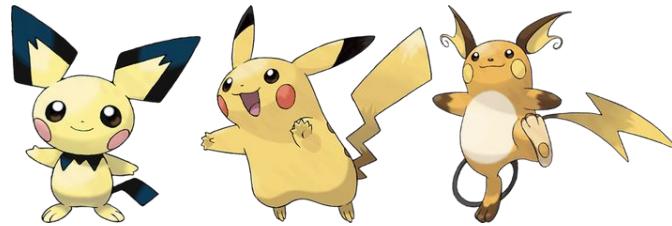
Supports structural evolution

Schema evolution changes are independent and free of side-effects, without rewriting files

Iceberg even supports partition evolution!!

```
CREATE TABLE orders_iceberg
  WITH (partitioning =
        ARRAY['month(orderdate)']) AS
SELECT * FROM tpch.sf1.orders;
```

```
ALTER TABLE orders_iceberg
  SET PROPERTIES partitioning =
    ARRAY['day(orderdate)');
```

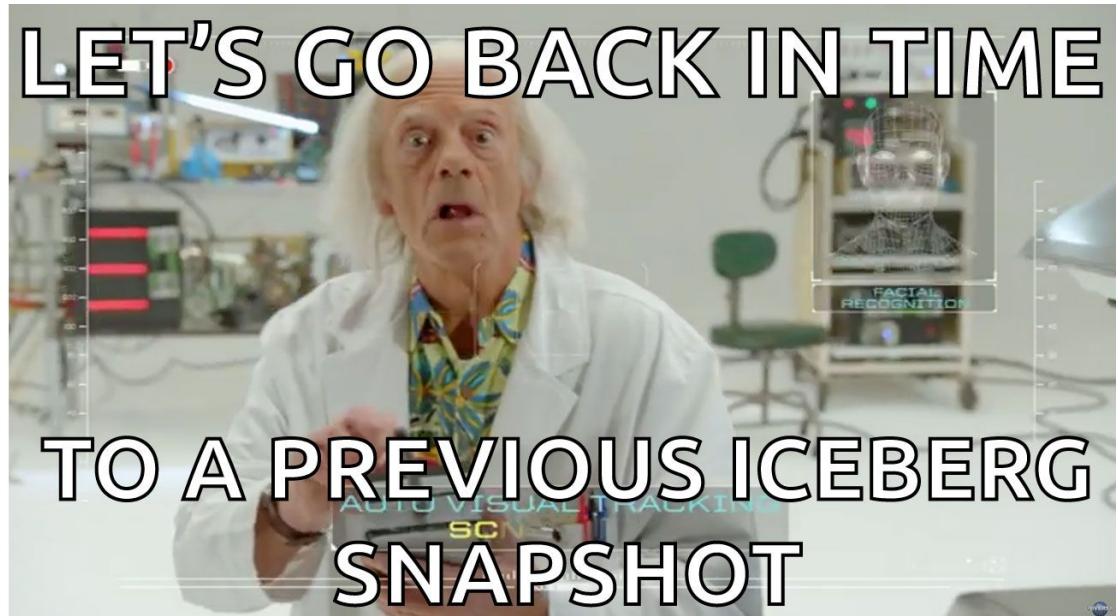


- orderdate_day=1996-06-13/
- orderdate_day=1996-09-26/
- orderdate_day=1997-11-04/
- orderdate_month=1992-01/
- orderdate_month=1992-02/
- orderdate_month=1992-03/
- orderdate_month=1992-04/
- orderdate_month=1992-05/

Table versioning via snapshots

Content and structural changes create new snapshots which enable

- Point in time reads (time-travel)
- Branching & tagging
- Rollbacks



Architecture overview

Metadata layer
lives in the
. /metadata dir

saved in the
. /data dir

Data layer



snapshots are
created anytime
data or structure is
changed

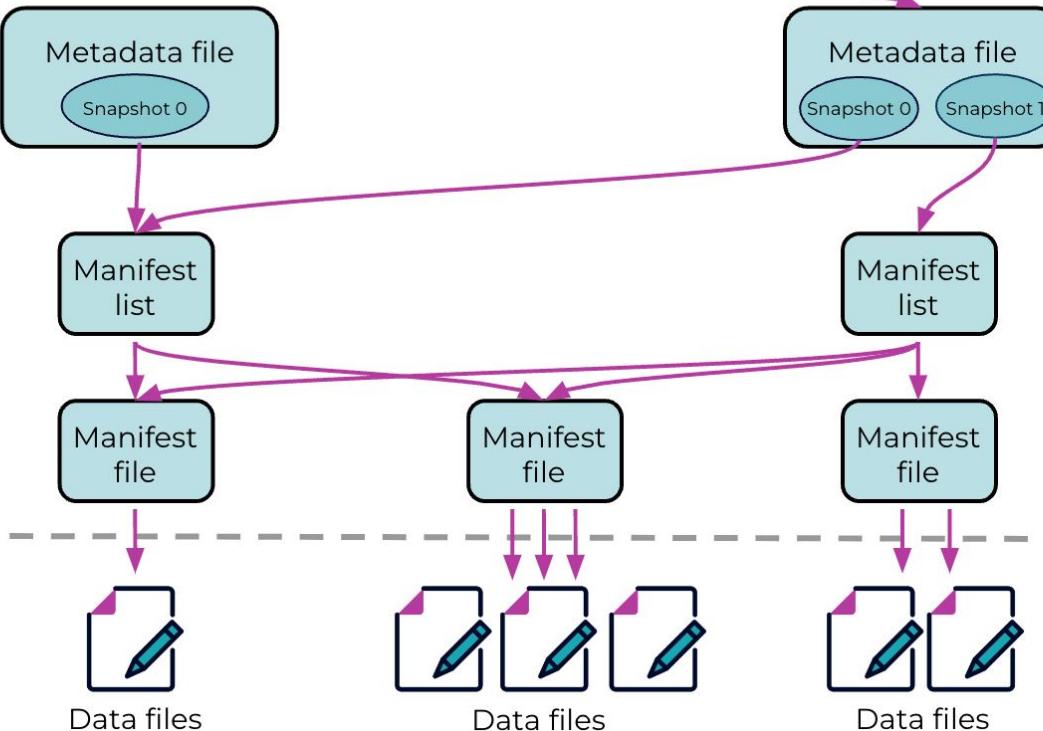
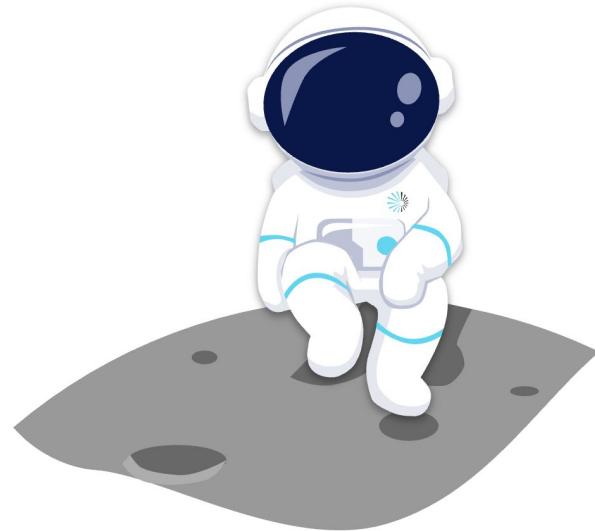


Table maintenance is required

Immutable files coupled with metadata sprawl are kept in check with cleanup functions

- Compaction
- Recalculate statistics
- Snapshot expiration
- Delete orphaned files



data maintenance jobs

Demo

Create, populate & query a table

Apache NiFi

What is it?

An easy to use,
powerful, and
reliable system to
process and
distribute data

NiFi automates cybersecurity, observability, event streams, and generative AI data pipelines and distribution for thousands of companies worldwide across every industry.

Download

View Documentation



Documentation

Development

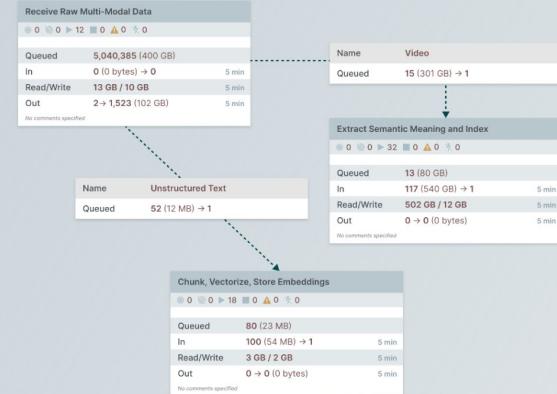
Community

Projects

Apache



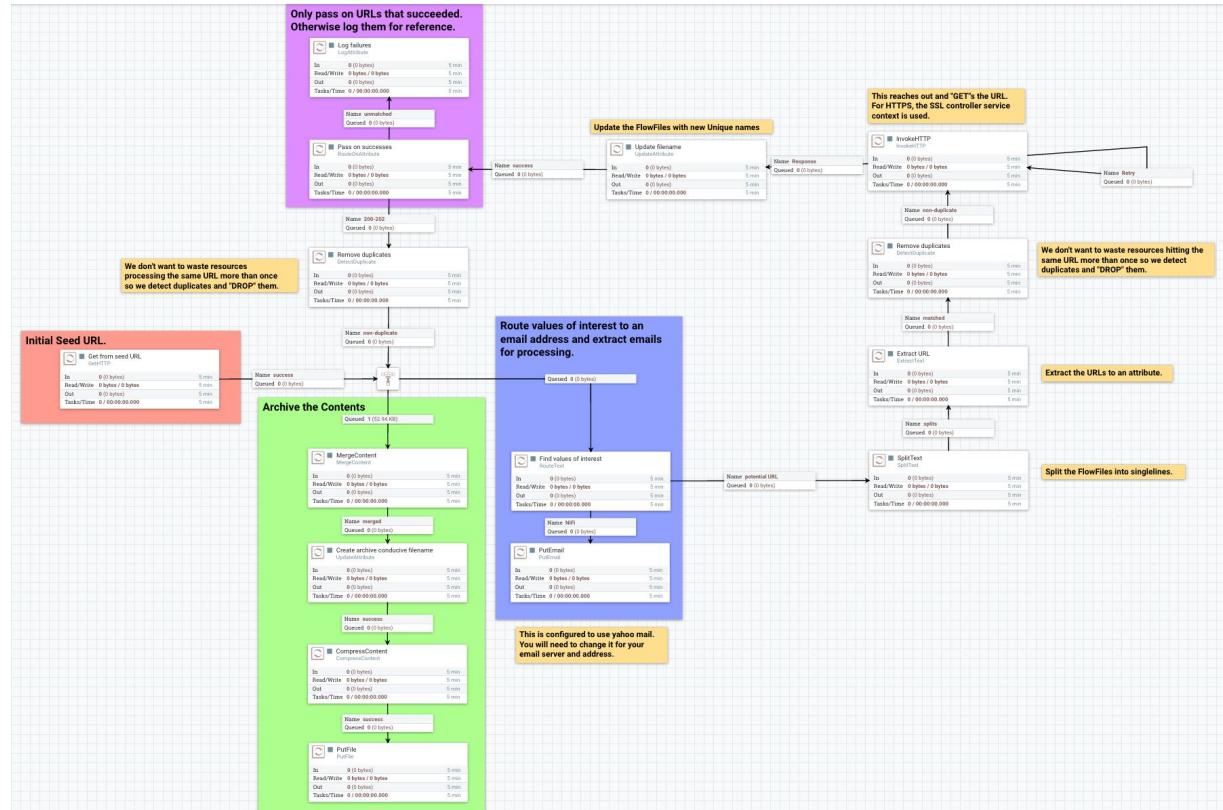
Download



Low/no-code dataflow (pipeline) creation tool

NiFi was built to **automate the flow of data between systems.**

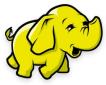
While the term '**dataflow**' is used in a variety of contexts, we use it here to mean the **automated and managed flow of information between systems.**



Check out my NiFi intro tutorial at <https://learn-with-lester.frebsite.online/tutorials/flow-dev-intro/>

Batch or streaming?

YES! Both!!



ICEBERG



splunk>



and many more connectors!

Development approach

 ListFile ListFile 2.0.0 org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 5 min
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min

 FetchFile FetchFile 2.0.0 org.apache.nifi - nifi-standard-nar
In 0 (0 bytes) 5 min
Read/Write 0 bytes / 0 bytes 
Out 0 (0 bytes) 5 min
Tasks/Time 0 / 00:00:00.000 5 min

Drag processors from the toolbar to the canvas

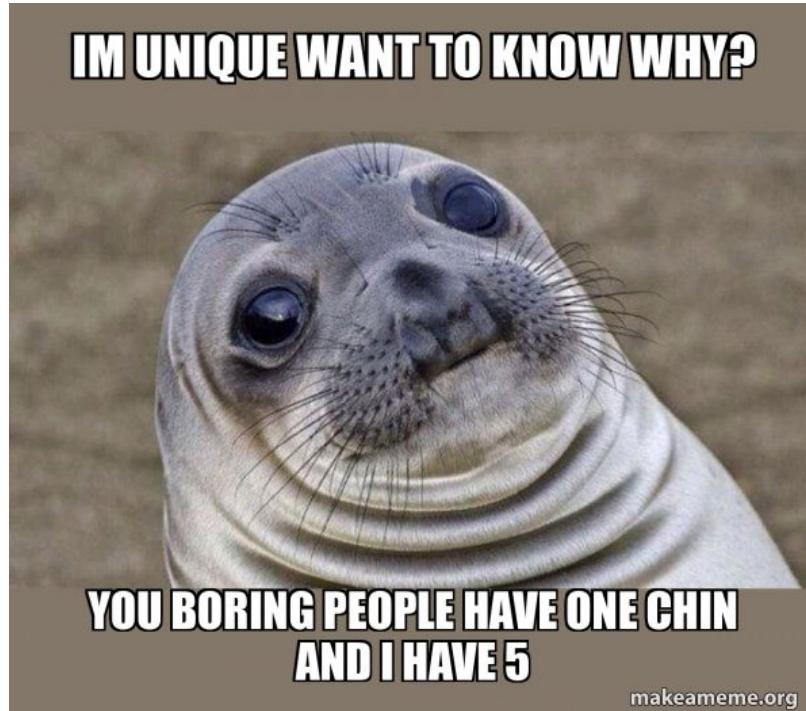
- Data ingest
- Transformations, modifications & conversions
- Activities such as table lookups
- Routing & mediation
- Data egress

Create connections between the processors

- Backed by queues
- Backpressure controls
- Configurable prioritization

NiFi tackles the full development life cycle

The NiFi UI also addresses the rest of the development life cycle



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

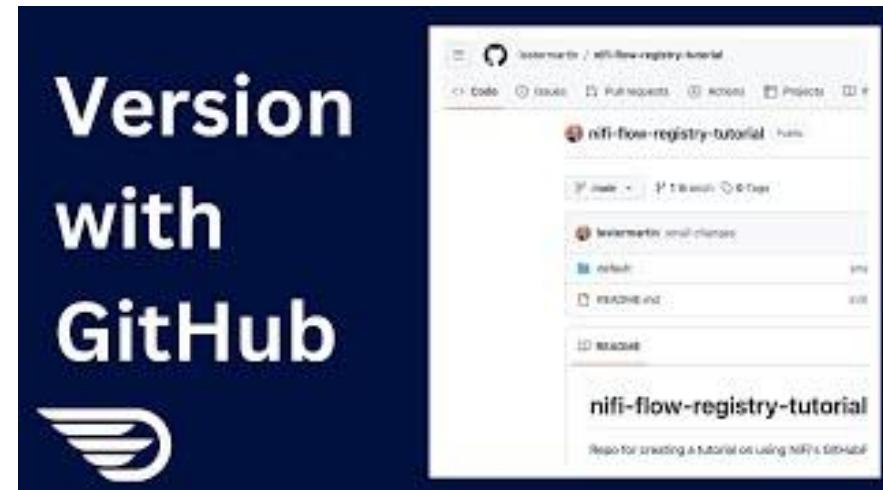
- **Compile & build** – No such thing as the dataflow is ready to run
 - Custom processors can be created with Python and Java



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

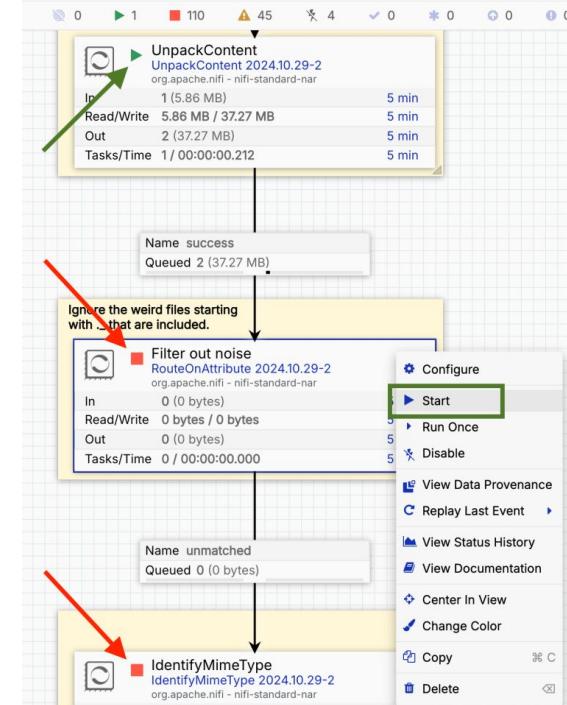
- **Compile & build** – No such thing as the dataflow is ready to run
 - Custom processors can be created with Python and Java
- **Deploy** – Just as above, nothing to do
 - Versioning and dev/test/prod movements possible &
recommended



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

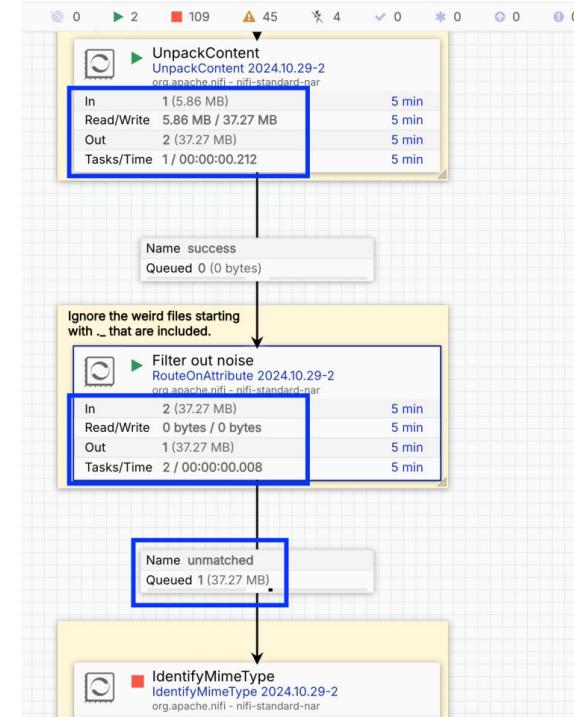
- **Compile & build** – No such thing as the dataflow is ready to run
 - Custom processors can be created with Python and Java
- **Deploy** – Just as above, nothing to do
 - Versioning and dev/test/prod movements possible & *recommended*
- **Runtime** – NiFi is the execution environment
 - Each processor is independently started/stopped
 - Can scale from a single node to 1000+ node clusters



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

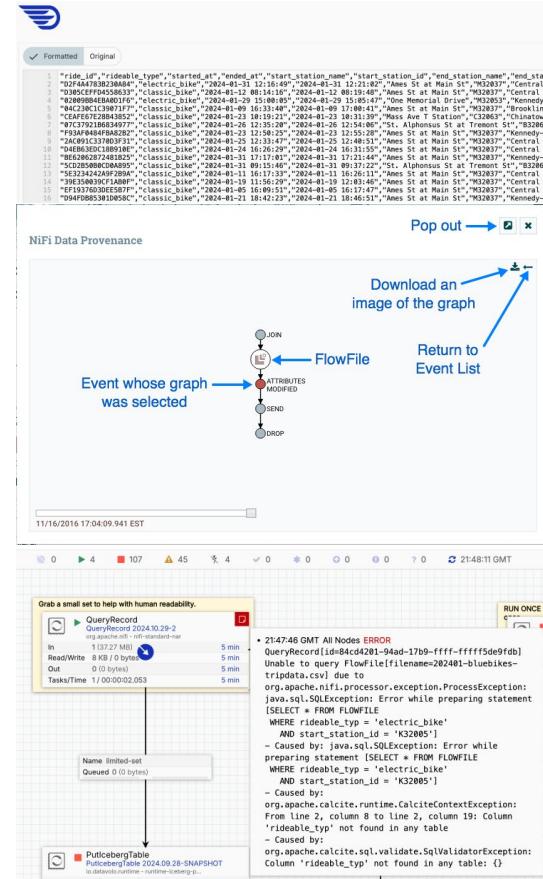
- **Compile & build** – No such thing as the dataflow is ready to run
 - Custom processors can be created with Python and Java
- **Deploy** – Just as above, nothing to do
 - Versioning and dev/test/prod movements possible & *recommended*
- **Runtime** – NiFi is the execution environment
 - Each processor is independently started/stopped
 - Can scale from a single node to 1000+ node clusters
- **Monitoring & reporting** – Processors & connection queues provide detailed statistics



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

- **Compile & build** – No such thing as the dataflow is ready to run
 - Custom processors can be created with Python and Java
- **Deploy** – Just as above, nothing to do
 - Versioning and dev/test/prod movements possible & recommended
- **Runtime** – NiFi is the execution environment
 - Each processor is independently started/stopped
 - Can scale from a single node to 1000+ node clusters
- **Monitoring & reporting** – Processors & connection queues provide detailed statistics
- **Tracing & debugging** – Comprehensive data viewing & lineage exists as well as log viewing exists



NiFi tackles the full development life cycle

The NiFi UI also **addresses the rest of the development life cycle**

- **Compile & build** – No such thing as
○ Custom processors can be created
- **Deploy** – Just as above, nothing to do
○ Versioning and dev/test/prod recommended
- **Runtime** – NiFi is the execution environment
○ Each processor is independent
○ Can scale from a single node to thousands
- **Monitoring & reporting** – Processors provide detailed statistics
- **Tracing & debugging** – Comprehensive tracing exists as well as log viewing exists

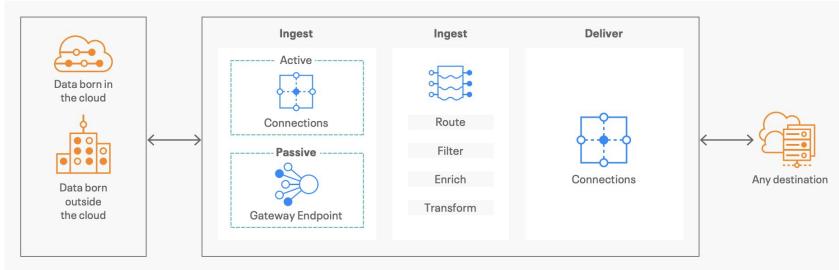


Demo

Explore a simple dataflow

Cloudera & Snowflake OSS+ products

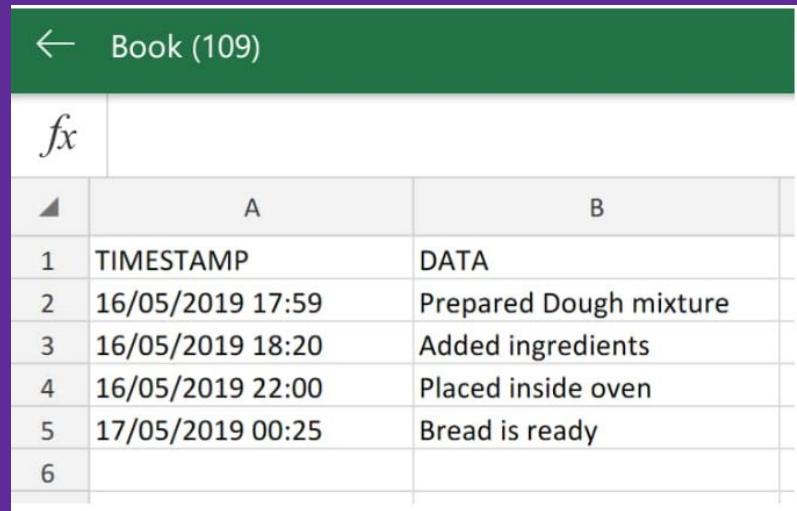
CLOUDERA DATAFLOW



powered by
APACHE **nifi**

Ingestion strategies

Timestamped immutable data



The screenshot shows a Microsoft Excel spreadsheet with a green header bar containing the text "← Book (109)". Below the header is a toolbar with icons for file operations. The main area contains a table with two columns, A and B. Column A is labeled "TIMESTAMP" and column B is labeled "DATA". The data rows are as follows:

	A	B
1	TIMESTAMP	DATA
2	16/05/2019 17:59	Prepared Dough mixture
3	16/05/2019 18:20	Added ingredients
4	16/05/2019 22:00	Placed inside oven
5	17/05/2019 00:25	Bread is ready
6		

Isn't there a PutIceberg processor?

Well, yes and no, but first here is what it does

- Bulks up a bunch of records or files coming into it
- Converts them into files using the Iceberg table's file format (ex: Apache Parquet)
- Moves those files to the table's location on the data lake
- Uses the Iceberg API to include the new records from the file(s) created (creates a new snapshot)

Sounds like a good thing, but...

Status of the PutIceberg processor

Cloudera Dataflow

Really only works with Cloudera Data Platform (i.e. Hadoop) which does allow object store persistence

Snowflake Openflow

A replacement processor, PutIcebergTable, was created, but only works with the Apache Polaris metadata catalog

Status of the PutIceberg processor

Cloudera Dataflow

Really only works with Cloudera Data Platform (i.e. Hadoop) which does allow object store persistence

Not OSS

Snowflake Openflow

A replacement processor, PutIcebergTable, was created, but only works with the Apache Polaris metadata catalog

Not OSS



Status of the PutIceberg processor

Cloudera Dataflow

Really only works with Cloudera Data Platform (i.e. Hadoop) which does allow object store persistence

Not OSS

Snowflake Openflow

A replacement processor, PutIcebergTable, was created, but only works with the Apache Polaris metadata catalog

Not OSS

Pure OSS Apache NiFi approach

Essentially, we just need to replicate the steps listed on prior slide

Strategically, “someone” needs to enhance PutIceberg to work with HDFS
and popular object stores & a variety of catalog/metastores (not just HMS)

Demo

Emulate Putlceberg behavior

Caveats of the PutIceberg strategy

Seems about perfect for append-only use cases, but there are concerns

- Data engineer needs to know which file format is used
- Difficult task to make sure the files meet other table configurations
 - Partitioning
 - Sorting
 - Bucketing
 - Bloom filters
- File type conversions happen on NiFi cluster instead of analytics compute engine
- Files can be deleted over time as part of table maintenance

Okay... is there an alternative?

Ephemeral external table strategy

Build a highly reusable dataflow that

- Bulks up a bunch of records or files coming into it – *leave it in whatever format it arrives in*
- Save files to data lake
- Create a Hive external table
- Perform INSERT INTO (or MERGE) SQL statement on the compute engine – *let the analytics cluster handle the file conversion and the any advanced table configurations*
- Drop ephemeral table
- Move or delete the data lake files

Just as easy as before!

The screenshot shows the 'Edit Processor' configuration for an 'ExecuteSQL 2.4.0' processor. The processor is currently set to 'CREATE / ITAS / DROP'. The 'Properties' tab is selected. A 'Required field' message is displayed above the table. The table lists properties and their values:

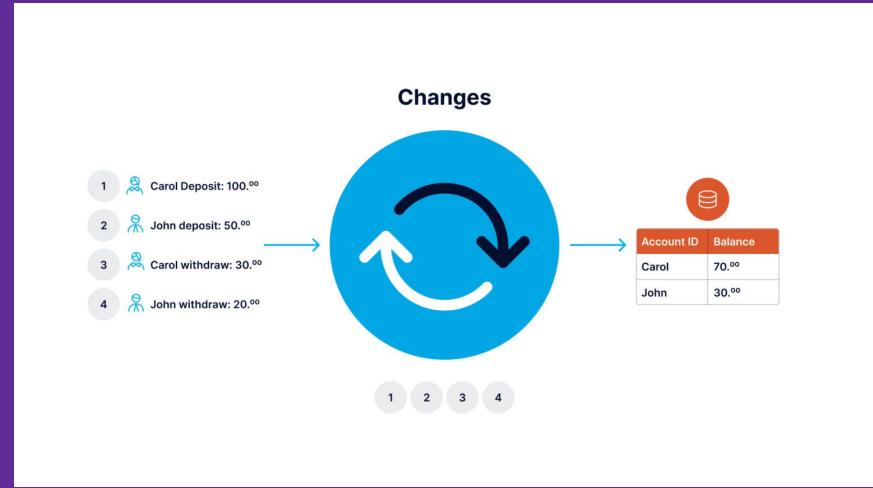
Property	Value
Database Connection Pooling Service	DBCPConnectionPool
SQL Pre-Query	CREATE TABLE mycloud.nifi_ingest.\$...
SQL Query	INSERT INTO mycloud.nifi_ingest.poke...
SQL Post-Query	DROP TABLE mycloud.nifi_ingest.\${uui...}

Demo

Use the ephemeral table strategy

Ingestion strategies

Table synchronization (CDC)



KISS principle can be appropriate

Keep it simple if

- The data really isn't BIG data
- You can wait a period of time for the sync-up
- You don't need to replay every change, you just need it to be accurate as-of refresh

KISS principle can be appropriate

Keep it simple if

- The data really isn't BIG data
- You can wait a period of time for the sync-up
- You don't need to replay every change, you just need it to be accurate as-of refresh

Build a simple pipeline that

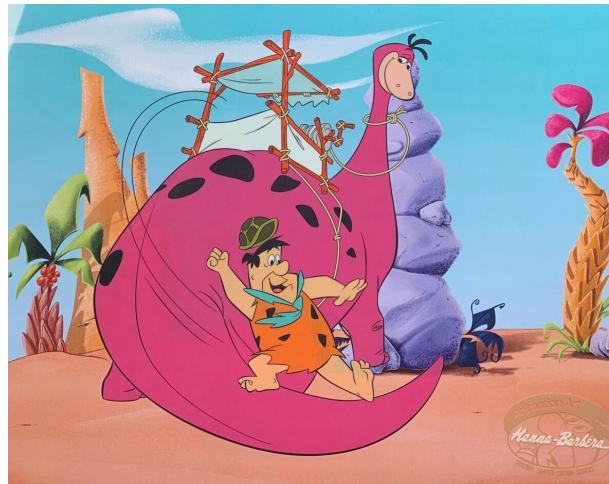
- Builds a new version of the table
- Drop (or rename) the active table
- Rename the replacement table to be the active table
- *Rinse, lather, repeat*



Otherwise... just use the ephemeral table strategy

Easy peasy...

Just swap out the INSERT statement with an appropriate MERGE statement!



Yabba
Dabba
Do!

Table maintenance

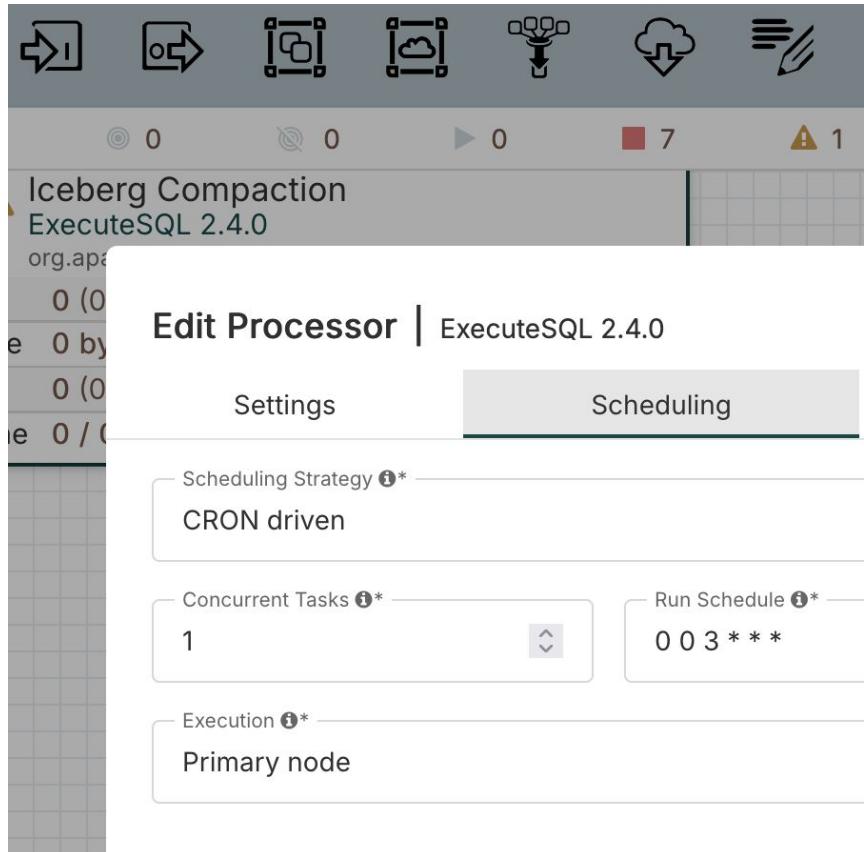
No problem...



Let NiFi be your scheduler and invoker

Processors can be run on a schedule which is perfect for

- Compaction
- Recalculate statistics
- Snapshot expiration
- Delete orphaned files



Thank you!

Lester Martin

<https://linktr.ee/lestermartin>

Trino Developer Advocate @ Starburst

<https://devcenter.starburst.io/>



Scan for a Trino and Iceberg cheat sheet



Source files - https://github.com/lestermartin/events/tree/main/2025-06-16_NiFi2Iceberg